

# GEN AI 인텐시브 과정

---

강사장철원

## Section 0

### 코스소개

DAY1

LLM  
Basic  
Concept

DAY2

Transformers  
paper  
review

DAY3

Transformers  
LangChain  
LangGraph

DAY4

DAY5

DAY6

LLM  
service  
develop

DAY7

Final Project

DAY8

## □ 실무적용

# GEN AI 인텐시브 과정

Section 1. 실무 적용

---

Section 1-1. 모델 성능

# 모델의 중요성

LLM 좋은 모델 사용 -> 성능 향상

## Section

### LangChain mcp adapter를 활용한 MCP 실습

# 라이브러리 설치

```
azureuser@b2b28-ML:~$ pyenv activate py3_11_9  
(py3_11_9) azureuser@b2b28-ML:~$ pip install langchain-mcp-adapters  
(py3_11_9) azureuser@b2b28-ML:~$ pip install nest_asyncio
```

## Section

### LangChain mcp adapter를 활용한 MCP 실습

# mcp\_server.py

```
(py3_11_9) azureuser@b2b28-ML:~/work/jupyter/3_llm_service$ vim mcp_server.py
```

## Section

### LangChain mcp adapter를 활용한 MCP 실습

# mcp\_server.py

```
import uuid
from typing import Dict
from mcp.server.fastmcp import FastMCP
from langchain_chroma import Chroma
from langchain_openai import AzureOpenAIEmbeddings, AzureChatOpenAI
from langchain_core.chat_history import InMemoryChatMessageHistory

AZURE_ENDPOINT = "https://b2b28-md6czaj-eastus2.cognitiveservices.azure.com/"
AZURE_API_KEY = "API키"
API_VERSION = "2024-12-01-preview"
EMBEDDING_DEPLOYMENT = "text-embedding-3-small"
CHAT_DEPLOYMENT = "gpt-4.1-mini"

azure_llm = AzureChatOpenAI(
    api_key=AZURE_API_KEY,
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=CHAT_DEPLOYMENT,
    openai_api_version=API_VERSION,
    temperature=0.3,
    max_tokens=800
)

embedding_model = AzureOpenAIEmbeddings(
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=EMBEDDING_DEPLOYMENT,
    api_key=AZURE_API_KEY,
    openai_api_version=API_VERSION
)

retriever = Chroma(
    persist_directory=".chromaDB",
    embedding_function=embedding_model
).as_retriever(search_kwargs={"k": 3})

chat_histories: Dict[str, InMemoryChatMessageHistory] = {}

def get_chat_history(session_id: str):
    if session_id not in chat_histories:
        chat_histories[session_id] = InMemoryChatMessageHistory()
    return chat_histories[session_id]

mcp = FastMCP("Agents")

@mcp.tool()
def rag_tool(question: str, session_id: str) -> Dict:
    chat_history = get_chat_history(session_id)
    history_text = "\n".join([m.content for m in chat_history.messages])
    docs = retriever.invoke(question)
    top_docs = docs[:3]
    context = "\n".join(doc.page_content for doc in top_docs)
    full_context = f"{history_text}\n\n{context}" if history_text else context

    prompt = f"\"{question}은 문맥입니다;\n{full_context}\n\n질문: {question}\n\n답변:\""
    result = azure_llm.invoke(prompt).content

    chat_history.add_user_message(question)
    chat_history.add_ai_message(result)

    return {
        "answer": result,
        "references": [doc.page_content for doc in top_docs]
    }

@mcp.tool()
def summarize_tool(text: str) -> str:
    prompt = f" 다음 텍스트를 한국어로 요약해줘:\n\n{text}"
    return azure_llm.invoke(prompt).content

@mcp.tool()
def rephrase_tool(text: str) -> str:
    prompt = f" 다음 문장을 정정하고 예의 바르게 바꿔줘:\n\n{text}"
    return azure_llm.invoke(prompt).content

if __name__ == "__main__":
    mcp.run(transport="stdio")
```

```
import uuid
from typing import Dict
from mcp.server.fastmcp import FastMCP
from langchain_chroma import Chroma
from langchain_openai import AzureOpenAIEmbeddings, AzureChatOpenAI
from langchain_core.chat_history import InMemoryChatMessageHistory

AZURE_ENDPOINT = "https://b2b28-md6czaj-eastus2.cognitiveservices.azure.com/"
AZURE_API_KEY = "API키"
API_VERSION = "2024-12-01-preview"
EMBEDDING_DEPLOYMENT = "text-embedding-3-small"
CHAT_DEPLOYMENT = "gpt-4.1-mini"

azure_llm = AzureChatOpenAI(
    api_key=AZURE_API_KEY,
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=CHAT_DEPLOYMENT,
    openai_api_version=API_VERSION,
    temperature=0.3,
    max_tokens=800
)

embedding_model = AzureOpenAIEmbeddings(
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=EMBEDDING_DEPLOYMENT,
    api_key=AZURE_API_KEY,
    openai_api_version=API_VERSION
)

retriever = Chroma(
    persist_directory=".chromaDB",
    embedding_function=embedding_model
).as_retriever(search_kwargs={"k": 3})

chat_histories: Dict[str, InMemoryChatMessageHistory] = {}

def get_chat_history(session_id: str):
    if session_id not in chat_histories:
        chat_histories[session_id] = InMemoryChatMessageHistory()
    return chat_histories[session_id]
```

## Section

### LangChain mcp adapter를 활용한 MCP 실습

# mcp\_server.py

```
import uuid
from typing import Dict
from mcp.server.fastmcp import FastMCP
from langchain_chroma import Chroma
from langchain_openai import AzureOpenAIEmbeddings, AzureChatOpenAI
from langchain_core.chat_history import InMemoryChatMessageHistory

AZURE_ENDPOINT = "https://b2b28-md6czaj-eastus2.cognitiveservices.azure.com/"
AZURE_API_KEY = "APt9I"
API_VERSION = "2024-12-01-preview"
EMBEDDING_DEPLOYMENT = "text-embedding-3-small"
CHAT_DEPLOYMENT = "gpt-4-1-mini"

azure_llm = AzureChatOpenAI(
    api_key=AZURE_API_KEY,
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=CHAT_DEPLOYMENT,
    openai_api_version=API_VERSION,
    temperature=0.3,
    max_tokens=800
)

embedding_model = AzureOpenAIEmbeddings(
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=EMBEDDING_DEPLOYMENT,
    api_key=AZURE_API_KEY,
    openai_api_version=API_VERSION
)

retriever = Chroma(
    persist_directory=".chromaDB",
    embedding_function=embedding_model
).as_retriever(search_kwargs={"k": 3})

chat_histories: Dict[str, InMemoryChatMessageHistory] = {}

def get_chat_history(session_id: str):
    if session_id not in chat_histories:
        chat_histories[session_id] = InMemoryChatMessageHistory()
    return chat_histories[session_id]

mcp = FastMCP("Agents")

@mcp.tool()
def rag_tool(question: str, session_id: str) -> Dict:
    chat_history = get_chat_history(session_id)
    history_text = "\n".join([m.content for m in chat_history.messages])
    docs = retriever.invoke(question)
    top_docs = docs[:3]
    context = "\n".join(doc.page_content for doc in top_docs)
    full_context = f"{history_text}\n\n{context}" if history_text else context
    prompt = f"""다음은 문맥입니다:\n{full_context}\n\n질문: {question}\n\n답변:"""
    result = azure_llm.invoke(prompt).content

    chat_history.add_user_message(question)
    chat_history.add_ai_message(result)

    return {
        "answer": result,
        "references": [doc.page_content for doc in top_docs]
    }

@mcp.tool()
def summarize_tool(text: str) -> str:
    prompt = f"다음 텍스트를 한국어로 요약해줘:\n\n{text}"
    return azure_llm.invoke(prompt).content

@mcp.tool()
def rephrase_tool(text: str) -> str:
    prompt = f"다음 문장을 정중하고 예의 바르게 바꿔줘:\n\n{text}"
    return azure_llm.invoke(prompt).content

if __name__ == "__main__":
    mcp.run(transport="stdio")
```

```
mcp = FastMCP("Agents")

@mcp.tool()
def rag_tool(question: str, session_id: str) -> Dict:
    chat_history = get_chat_history(session_id)
    history_text = "\n".join([m.content for m in chat_history.messages])
    docs = retriever.invoke(question)
    top_docs = docs[:3]
    context = "\n".join(doc.page_content for doc in top_docs)
    full_context = f"{history_text}\n\n{context}" if history_text else context
    prompt = f"""다음은 문맥입니다:\n{full_context}\n\n질문: {question}\n\n답변:"""
    result = azure_llm.invoke(prompt).content

    chat_history.add_user_message(question)
    chat_history.add_ai_message(result)

    return {
        "answer": result,
        "references": [doc.page_content for doc in top_docs]
    }

@mcp.tool()
def summarize_tool(text: str) -> str:
    prompt = f"다음 텍스트를 한국어로 요약해줘:\n\n{text}"
    return azure_llm.invoke(prompt).content

@mcp.tool()
def rephrase_tool(text: str) -> str:
    prompt = f"다음 문장을 정중하고 예의 바르게 바꿔줘:\n\n{text}"
    return azure_llm.invoke(prompt).content

if __name__ == "__main__":
    mcp.run(transport="stdio")
```

## Section

### LangChain mcp adapter를 활용한 MCP 실습

# mcp\_client.py

```
(py3_11_9) azureuser@b2b28-ML:~/work/jupyter/3_llm_service$ vim mcp_client.py
```

## Section

### LangChain mcp adapter를 활용한 MCP 실습

# mcp\_client.py

```
import asyncio
import uuid
from mcp import ClientSession, StdioServerParameters
from mcp.client.stdio import stdio_client
from langchain_mcp_adapters.tools import load_mcp_tools
from langgraph.prebuilt import create_react_agent
from langchain_openai import AzureChatOpenAI

AZURE_ENDPOINT = "https://b2b28-md6zcjaj-eastus2.cognitiveservices.azure.com/"
AZURE_API_KEY = "APIKey"
API_VERSION = "2024-12-01-preview"
CHAT_DEPLOYMENT = "gpt-4.1-mini"

model = AzureChatOpenAI(
    api_key=AZURE_API_KEY,
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=CHAT_DEPLOYMENT,
    openai_api_version=API_VERSION,
    temperature=0.3,
)

async def main():
    session_id = str(uuid.uuid4())

    server_params = StdioServerParameters(
        command="python",
        args=["mcp_server.py"], # MCP 서버 경로
    )

    async with stdio_client(server_params) as (read, write):
        async with ClientSession(read, write) as session:
            await session.initialize()
            tools = await load_mcp_tools(session)
            agent = create_react_agent(model, tools)

            while True:
                question = input("\n질문을 입력해주세요: ").strip()
                if not question:
                    break

                if "오악" in question:
                    tool_prompt = f"Use summarize_tool to summarize:\n\n{question}"
                elif "정중" in question or "공손" in question or "예의" in question:
                    tool_prompt = f"Use rephrase_tool to rephrase:\n\n{question}"
                else:
                    tool_prompt = f"Use rag_tool to answer:\nquestion: {question}\nsession_id: {session_id}"

                response = await agent.invoke({"messages": tool_prompt})

                print("\n[에이전트 응답]")
                print(response)

                cont = input("\n계속하시겠습니까? (예/아니오): ").strip()
                if not cont.startswith("에"):
                    break

if __name__ == "__main__":
    asyncio.run(main())
```

```
import asyncio
import uuid
from mcp import ClientSession, StdioServerParameters
from mcp.client.stdio import stdio_client
from langchain_mcp_adapters.tools import load_mcp_tools
from langgraph.prebuilt import create_react_agent
from langchain_openai import AzureChatOpenAI

AZURE_ENDPOINT = "https://b2b28-md6zcjaj-eastus2.cognitiveservices.azure.com/"
AZURE_API_KEY = "APIKey"
API_VERSION = "2024-12-01-preview"
CHAT_DEPLOYMENT = "gpt-4.1-mini"

model = AzureChatOpenAI(
    api_key=AZURE_API_KEY,
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=CHAT_DEPLOYMENT,
    openai_api_version=API_VERSION,
    temperature=0.3,
)
```

## Section

### LangChain mcp adapter를 활용한 MCP 실습

# mcp\_client.py

```
import asyncio
import uuid
from mcp import ClientSession, StdioServerParameters
from mcp.client.stdio import stdio_client
from langchain_mcp_adapters.tools import load_mcp_tools
from langgraph.prebuilt import create_react_agent
from langchain_openai import AzureChatOpenAI

AZURE_ENDPOINT = "https://b2b28-md6zczaj-eastus2.cognitiveservices.azure.com/"
AZURE_API_KEY = "APIKey"
API_VERSION = "2024-12-01-preview"
CHAT_DEPLOYMENT = "gpt-4.1-mini"

model = AzureChatOpenAI(
    api_key=AZURE_API_KEY,
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=CHAT_DEPLOYMENT,
    openai_api_version=API_VERSION,
    temperature=0.3,
)

async def main():
    session_id = str(uuid.uuid4())

    server_params = StdioServerParameters(
        command="python",
        args=["mcp_server.py"], # MCP 서버 경로
    )

    async with stdio_client(server_params) as (read, write):
        async with ClientSession(read, write) as session:
            await session.initialize()
            tools = await load_mcp_tools(session)
            agent = create_react_agent(model, tools)

            while True:
                question = input("\n질문을 입력해주세요: ").strip()
                if not question:
                    break

                if "요약" in question:
                    tool_prompt = f"Use summarize_tool to summarize:\n\n{question}"
                elif "정중" in question or "공손" in question or "예의" in question:
                    tool_prompt = f"Use rephrase_tool to rephrase:\n\n{question}"
                else:
                    tool_prompt = f"Use rag_tool to answer:\nquestion: {question}\nsession_id: {session_id}"

                response = await agent.ainvoke({"messages": tool_prompt})

                print("\n[에이전트 응답]")
                print(response)

                cont = input("\n계속하시겠습니까? (예/아니오): ").strip()
                if not cont.startswith("예"):
                    break

    if __name__ == "__main__":
        asyncio.run(main())
```

```
async def main():
    session_id = str(uuid.uuid4())

    server_params = StdioServerParameters(
        command="python",
        args=["mcp_server.py"], # MCP 서버 경로
    )

    async with stdio_client(server_params) as (read, write):
        async with ClientSession(read, write) as session:
            await session.initialize()
            tools = await load_mcp_tools(session)
            agent = create_react_agent(model, tools)

            while True:
                question = input("\n질문을 입력해주세요: ").strip()
                if not question:
                    break

                if "요약" in question:
                    tool_prompt = f"Use summarize_tool to summarize:\n\n{question}"
                elif "정중" in question or "공손" in question or "예의" in question:
                    tool_prompt = f"Use rephrase_tool to rephrase:\n\n{question}"
                else:
                    tool_prompt = f"Use rag_tool to answer:\nquestion: {question}\nsession_id: {session_id}"

                response = await agent.ainvoke({"messages": tool_prompt})

                print("\n[에이전트 응답]")
                print(response)

                cont = input("\n계속하시겠습니까? (예/아니오): ").strip()
                if not cont.startswith("예"):
                    break

    if __name__ == "__main__":
        asyncio.run(main())
```

## Section

### LangChain mcp adapter를 활용한 MCP 실습

## 실행

```
(py3_11_9) azureuser@b2b28-ML:~/work/jupyter/3_llm_service$ python mcp_client.py
[07/21/25 03:59:20] INFO      Processing request of type ListToolsRequest

질문을 입력해 주세요 : 다음 문장 공손하게 바꿔 줘. "이 메일 빨리 보내세요"
[07/21/25 03:59:35] INFO      Processing request of type CallToolRequest
[07/21/25 03:59:36] INFO      HTTP Request: POST
https://b2b28-md6zcjaj-eastus2.cognitiveservices.azure.com/openai/deployments/gpt-4.1-mini/chat/completions?api-version=2024-12-01-preview "HTTP/1.1 200 OK"

[에이전트 응답]
{
  "messages": [
    {
      "role": "HumanMessage",
      "content": "Use rephrase_tool to rephrase:\n\n다음 문장 공손하게 바꿔 줘. \"이 메일 빨리 보내세요\""
    },
    {
      "role": "AIMessage",
      "content": "시일 내에 보내주시면 감사하겠습니다.",
      "additional_kwargs": {}
    }
  ]
}
```

## Section

### LangChain mcp adapter를 활용한 MCP 실습

## 실행

```
계 속 하 시 겠 습 니 까 ? (예 /아 니 오 ) : 예
질문 을 입 력 해 주 세 요 : ARS 시스 템 점 검 은 언 제 하 나요 ?
[07/21/25 04:00:05] INFO Processing request of type CallToolRequest server.py:625
[07/21/25 04:00:06] INFO HTTP Request: POST client.py:1025
https://b2b28-md6czcraj-eastus2.cognitiveservices.azure.com/openai/deployments/text-embedding-3-small/embeddings?api-version=2024-12-01-preview "HTTP/1.1 200 OK"
[07/21/25 04:00:07] INFO HTTP Request: POST client.py:1025
https://b2b28-md6czcraj-eastus2.cognitiveservices.azure.com/openai/deployments/gpt-4.1-mini/chat/completions?api-version=2024-12-01-preview "HTTP/1.1 200 OK"

[에 이 전 트 응답 ]
{'messages': [HumanMessage(content='Use rag_tool to answer:\nquestion: ARS 시스 템 점 검 은 언 제 하 나요?\nsession_id: f954ca2f-53f4-4bb0-a9bc-176ab5b68366', additional_kwargs={}, response_metadata={}, id='57c99a6b-4caa-453e-ae66-4708e1156800'), AIMessage(content='', additional_kwargs={'tool_calls': [{'id': 'call_4KXsqmV69w9rBjFWdDZR9ANQ', 'function': {'arguments': {'question': 'ARS 시스 템 점 검 은 언 제 하 나요?'}, 'session_id': 'f954ca2f-53f4-4bb0-a9bc-176ab5b68366'}, 'name': 'rag_tool', 'type': 'function'}], 'refusal': None}, response_metadata={'token_usage': {'completion_tokens': 51, 'prompt_tokens': 115, 'total_tokens': 166, 'completion_tokens_details': {'accepted_prediction_tokens': 0, 'audio_tokens': 0, 'reasoning_tokens': 0, 'rejected_prediction_tokens': 0}, 'prompt_tokens_details': {'audio_tokens': 0, 'cached_tokens': 0}}, 'model_name': 'gpt-4.1-mini-2025-04-14', 'system_fingerprint': 'fp_178c8d546f', 'id': 'chatcmpl-Bvbz2FgTEzbDsYp4XkT1FqQY6T1Mk', 'service_tier': None, 'prompt_filter_results': [{'prompt_index': 0, 'content_filter_results': {'hate': {'filtered': False, 'severity': 'safe'}, 'self_harm': {'filtered': False, 'severity': 'safe'}, 'sexual': {'filtered': False, 'severity': 'safe'}, 'violence': {'filtered': False, 'severity': 'safe'}}}], 'finish_reason': 'tool_calls', 'logprobs': None, 'content_filter_results': {}, 'id': 'run--59392d48-5461-4cc8-931e-1932879caa40-0', 'tool_calls':[{'name': 'rag_tool', 'args': {'question': 'ARS 시스 템 점 검 은 언 제 하 나요?'}, 'session_id': 'f954ca2f-53f4-4bb0-a9bc-176ab5b68366'}, {'id': 'call_4KXsqmV69w9rBjFWdDZR9ANQ', 'type': 'tool_call'}], 'usage_metadata':{'input_tokens': 115, 'output_tokens': 51, 'total_tokens': 166, 'input_token_details': {'audio': 0, 'cache_read': 0}, 'output_token_details': {'audio': 0, 'reasoning': 0}}}, ToolMessage(content='\n    "answer": "ARS 시스 템 점 검 은 매 월 수 행 합 니다.",\n    "references": [\n        "ARS 시스 템은 매 월 점 검 하며 장애 발 생 시 즉 시 전 산 팀에 보 고 합 니다.",\n        "기 지 국 점 검 은 월 1회 정 기 적 으로 수 행 되 며 장애 이 력은 3년 간 보 관 합 니다.",\n        "고객 접 점 에서 확 인 된 오 류는 VOC 시스 템을 통 해 수 집 및 분 류 됩 니 다.\n    ]\n}', name='rag_tool', id='0a2e0422-65t2-404c-90e1-5352f2244b20', tool_call_id='call_4KXsqmV69w9rBjFWdDZR9ANQ'), AIMessage(content='ARS 시스 템 점 검 은 매 월 수 행 합 니다. 점 검 시 장애 가 발 생 하면 즉 시 전 산 팀에 보 고 됩 니 다.', additional_kwargs={'refusal': None}, response_metadata={'token_usage': {'completion_tokens': 28, 'prompt_tokens': 271, 'total_tokens': 299, 'completion_tokens_details': {'accepted_prediction_tokens': 0, 'audio_tokens': 0, 'reasoning_tokens': 0, 'rejected_prediction_tokens': 0}, 'prompt_tokens_details': {'audio_tokens': 0, 'cached_tokens': 0}}, 'model_name': 'gpt-4.1-mini-2025-04-14', 'system_fingerprint': 'fp_178c8d546f', 'id': 'chatcmpl-Bvbz5ql8MV6bfJptxszMMPvCvN0B', 'service_tier': None, 'prompt_filter_results': [{'prompt_index': 0, 'content_filter_results': {'hate': {'filtered': False, 'severity': 'safe'}, 'self_harm': {'filtered': False, 'severity': 'safe'}, 'sexual': {'filtered': False, 'severity': 'safe'}, 'violence': {'filtered': False, 'severity': 'safe'}}}], 'finish_reason': 'stop', 'logprobs': None, 'content_filter_results': {'hate': {'filtered': False, 'severity': 'safe'}, 'self_harm': {'filtered': False, 'severity': 'safe'}, 'sexual': {'filtered': False, 'severity': 'safe'}, 'violence': {'filtered': False, 'severity': 'safe'}}, 'id': 'run--85c4d542-7af4-4d51-aecc-464791822e27-0', 'usage_metadata':{'input_tokens': 271, 'output_tokens': 28, 'total_tokens': 299, 'input_token_details': {'audio': 0, 'cache_read': 0}, 'output_token_details': {'audio': 0, 'reasoning': 0}}})]

계 속 하 시 겠 습 니 까 ? (예 /아 니 오 ) : 아 니 오
(py3_11_9) azureuser@b2b28-ML:~/work/jupyter/3_llm_service$
```

# MCP

Section 2. MCP 실습

---

**Section 2-3. ReAct Agent 기반 실습**

# ReAct

## Reasoning + Acting

질문의 의미를 분석하고,  
어떤 도구를 사용할지 결정

자신이 결정한  
도구를 사용

## Section

### ReAct Agent 기반 실습

# mcp\_server.py(변동X)

```
import uuid
from typing import Dict
from mcp.server.fastmcp import FastMCP
from langchain_chroma import Chroma
from langchain_openai import AzureOpenAIEmbeddings, AzureChatOpenAI
from langchain_core.chat_history import InMemoryChatMessageHistory

AZURE_ENDPOINT = "https://b2b28-md6czaj-eastus2.cognitiveservices.azure.com/"
AZURE_API_KEY = "API키"
API_VERSION = "2024-12-01-preview"
EMBEDDING_DEPLOYMENT = "text-embedding-3-small"
CHAT_DEPLOYMENT = "gpt-4.1-mini"

azure_llm = AzureChatOpenAI(
    api_key=AZURE_API_KEY,
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=CHAT_DEPLOYMENT,
    openai_api_version=API_VERSION,
    temperature=0.3,
    max_tokens=800
)

embedding_model = AzureOpenAIEmbeddings(
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=EMBEDDING_DEPLOYMENT,
    api_key=AZURE_API_KEY,
    openai_api_version=API_VERSION
)

retriever = Chroma(
    persist_directory=".chromaDB",
    embedding_function=embedding_model
).as_retriever(search_kwargs={"k": 3})

chat_histories: Dict[str, InMemoryChatMessageHistory] = {}

def get_chat_history(session_id: str):
    if session_id not in chat_histories:
        chat_histories[session_id] = InMemoryChatMessageHistory()
    return chat_histories[session_id]

mcp = FastMCP("Agents")

@mcp.tool()
def rag_tool(question: str, session_id: str) -> Dict:
    chat_history = get_chat_history(session_id)
    history_text = "\n".join([m.content for m in chat_history.messages])
    docs = retriever.invoke(question)
    top_docs = docs[:3]
    context = "\n".join(doc.page_content for doc in top_docs)
    full_context = f"{history_text}\n\n{context}" if history_text else context
    prompt = f"\"{question}을 문맥입니다;\n\n{full_context}\n\n질문: {question}\n\n답변:\""
    result = azure_llm.invoke(prompt).content

    chat_history.add_user_message(question)
    chat_history.add_ai_message(result)

    return {
        "answer": result,
        "references": [doc.page_content for doc in top_docs]
    }

@mcp.tool()
def summarize_tool(text: str) -> str:
    prompt = f" 다음 텍스트를 한국어로 요약해줘:\n\n{text}"
    return azure_llm.invoke(prompt).content

@mcp.tool()
def rephrase_tool(text: str) -> str:
    prompt = f" 다음 문장을 정정하고 예의 바르게 바꿔줘:\n\n{text}"
    return azure_llm.invoke(prompt).content

if __name__ == "__main__":
    mcp.run(transport="stdio")
```

```
import uuid
from typing import Dict
from mcp.server.fastmcp import FastMCP
from langchain_chroma import Chroma
from langchain_openai import AzureOpenAIEmbeddings, AzureChatOpenAI
from langchain_core.chat_history import InMemoryChatMessageHistory

AZURE_ENDPOINT = "https://b2b28-md6czaj-eastus2.cognitiveservices.azure.com/"
AZURE_API_KEY = "API키"
API_VERSION = "2024-12-01-preview"
EMBEDDING_DEPLOYMENT = "text-embedding-3-small"
CHAT_DEPLOYMENT = "gpt-4.1-mini"

azure_llm = AzureChatOpenAI(
    api_key=AZURE_API_KEY,
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=CHAT_DEPLOYMENT,
    openai_api_version=API_VERSION,
    temperature=0.3,
    max_tokens=800
)

embedding_model = AzureOpenAIEmbeddings(
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=EMBEDDING_DEPLOYMENT,
    api_key=AZURE_API_KEY,
    openai_api_version=API_VERSION
)

retriever = Chroma(
    persist_directory=".chromaDB",
    embedding_function=embedding_model
).as_retriever(search_kwargs={"k": 3})

chat_histories: Dict[str, InMemoryChatMessageHistory] = {}

def get_chat_history(session_id: str):
    if session_id not in chat_histories:
        chat_histories[session_id] = InMemoryChatMessageHistory()
    return chat_histories[session_id]
```

## Section

### ReAct Agent 기반 실습

# mcp\_server.py(변동X)

```
import uuid
from typing import Dict
from mcp.server.fastmcp import FastMCP
from langchain_chroma import Chroma
from langchain_openai import AzureOpenAIEmbeddings, AzureChatOpenAI
from langchain_core.chat_history import InMemoryChatMessageHistory

AZURE_ENDPOINT = "https://b2b28-md6czaj-eastus2.cognitiveservices.azure.com/"
AZURE_API_KEY = "APt9I"
API_VERSION = "2024-12-01-preview"
EMBEDDING_DEPLOYMENT = "text-embedding-3-small"
CHAT_DEPLOYMENT = "gpt-4-1-min"
AZURE_CHAT_DEPLOYMENT = "gpt-4-1-min"

azure_llm = AzureChatOpenAI(
    api_key=AZURE_API_KEY,
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=CHAT_DEPLOYMENT,
    openai_api_version=API_VERSION,
    temperature=0.3,
    max_tokens=800
)

embedding_model = AzureOpenAIEmbeddings(
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=EMBEDDING_DEPLOYMENT,
    api_key=AZURE_API_KEY,
    openai_api_version=API_VERSION
)

retriever = Chroma(
    persist_directory=".chromaDB",
    embedding_function=embedding_model
).as_retriever(search_kwargs={"k": 3})

chat_histories: Dict[str, InMemoryChatMessageHistory] = {}

def get_chat_history(session_id: str):
    if session_id not in chat_histories:
        chat_histories[session_id] = InMemoryChatMessageHistory()
    return chat_histories[session_id]

mcp = FastMCP("Agents")

@mcp.tool()
def rag_tool(question: str, session_id: str) -> Dict:
    chat_history = get_chat_history(session_id)
    history_text = "\n".join([m.content for m in chat_history.messages])
    docs = retriever.invoke(question)
    top_docs = docs[:3]
    context = "\n".join(doc.page_content for doc in top_docs)
    full_context = f"{history_text}\n\n{context}" if history_text else context

    prompt = f""" 다음은 문맥입니다:\n{full_context}\n\n질문: {question}\n\n답변: """
    result = azure_llm.invoke(prompt).content

    chat_history.add_user_message(question)
    chat_history.add_ai_message(result)

    return {
        "answer": result,
        "references": [doc.page_content for doc in top_docs]
    }

@mcp.tool()
def summarize_tool(text: str) -> str:
    prompt = f"다음 텍스트를 한국어로 요약해줘:\n\n{text}"
    return azure_llm.invoke(prompt).content

@mcp.tool()
def rephrase_tool(text: str) -> str:
    prompt = f"다음 문장을 정중하고 예의 바르게 바꿔줘:\n\n{text}"
    return azure_llm.invoke(prompt).content

if __name__ == "__main__":
    mcp.run(transport="stdio")
```

```
mcp = FastMCP("Agents")

@mcp.tool()
def rag_tool(question: str, session_id: str) -> Dict:
    chat_history = get_chat_history(session_id)
    history_text = "\n".join([m.content for m in chat_history.messages])
    docs = retriever.invoke(question)
    top_docs = docs[:3]
    context = "\n".join(doc.page_content for doc in top_docs)
    full_context = f"{history_text}\n\n{context}" if history_text else context

    prompt = f""" 다음은 문맥입니다:\n{full_context}\n\n질문: {question}\n\n답변: """
    result = azure_llm.invoke(prompt).content

    chat_history.add_user_message(question)
    chat_history.add_ai_message(result)

    return {
        "answer": result,
        "references": [doc.page_content for doc in top_docs]
    }

@mcp.tool()
def summarize_tool(text: str) -> str:
    prompt = f"다음 텍스트를 한국어로 요약해줘:\n\n{text}"
    return azure_llm.invoke(prompt).content

@mcp.tool()
def rephrase_tool(text: str) -> str:
    prompt = f"다음 문장을 정중하고 예의 바르게 바꿔줘:\n\n{text}"
    return azure_llm.invoke(prompt).content

if __name__ == "__main__":
    mcp.run(transport="stdio")
```

## Section

### ReAct Agent 기반 실습

# mcp\_client.py

```
(py3_11_9) azureuser@b2b28-ML:~/work/jupyter/3_llm_service$ vim mcp_react_client.py
```

## Section

### ReAct Agent 기반 실습

# mcp\_react\_client.py

```
import asyncio
import uuid
from mcp import ClientSession, StdioServerParameters
from mcp.client.stdio import stdio_client
from langchain_mcp_adapters.tools import load_mcp_tools
from langgraph.prebuilt import create_react_agent
from langchain_openai import AzureChatOpenAI

# Azure 설정
AZURE_ENDPOINT = "https://b2b28-md6czaj-eastus2.cognitiveservices.azure.com/"
AZURE_API_KEY = "APIKey"
API_VERSION = "2024-12-01-preview"
CHAT_DEPLOYMENT = "gpt-4.1-mini"

model = AzureChatOpenAI(
    api_key=AZURE_API_KEY,
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=CHAT_DEPLOYMENT,
    openai_api_version=API_VERSION,
    temperature=0.3,
    max_tokens=800
)

async def main():
    session_id = str(uuid.uuid4())

    # MCP 서버 실행 파라미터
    server_params = StdioServerParameters(
        command="python",
        args=["mcp_server.py"], # MCP 를 서버 경로
    )

    async with stdio_client(server_params) as (read, write):
        async with ClientSession(read, write) as session:
            await session.initialize()
            tools = await load_mcp_tools(session)

            # React Agent 생성 (툴 목록 제공)
            agent = create_react_agent(model, tools)

            while True:
                user_input = input("\n질문을 입력해주세요: ").strip()
                if not user_input:
                    break

                # React Agent에게 메시지 전달
                messages = [
                    {"role": "user", "content": f'{user_input}\n(session_id: {session_id})'}
                ]
                result = await agent.ainvoke({"messages": messages})

                print("\n[에이전트 응답]")
                print(result)

                cont = input("\n계속하시겠습니까? (예/아니오): ").strip()
                if not cont.startswith("예"):
                    break

if __name__ == "__main__":
    asyncio.run(main())
```

```
import asyncio
import uuid
from mcp import ClientSession, StdioServerParameters
from mcp.client.stdio import stdio_client
from langchain_mcp_adapters.tools import load_mcp_tools
from langgraph.prebuilt import create_react_agent
from langchain_openai import AzureChatOpenAI

# Azure 설정
AZURE_ENDPOINT = "https://b2b28-md6czaj-eastus2.cognitiveservices.azure.com/"
AZURE_API_KEY = "APIKey"
API_VERSION = "2024-12-01-preview"
CHAT_DEPLOYMENT = "gpt-4.1-mini"

model = AzureChatOpenAI(
    api_key=AZURE_API_KEY,
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=CHAT_DEPLOYMENT,
    openai_api_version=API_VERSION,
    temperature=0.3,
    max_tokens=800
)
```

## Section

### ReAct Agent 기반 실습

# mcp\_react\_client.py

```
import asyncio
import uuid
from mcp import ClientSession, StdioServerParameters
from mcp.client.stdio import stdio_client
from langchain_mcp_adapters.tools import load_mcp_tools
from langgraph.prebuilt import create_react_agent
from langchain_openai import AzureChatOpenAI

# Azure 설정
AZURE_ENDPOINT = "https://b2b28-md6czaj-eastus2.cognitiveservices.azure.com/"
AZURE_API_KEY = "APIKey"
API_VERSION = "2024-12-01-preview"
CHAT_DEPLOYMENT = "gpt-4.1-mini"

model = AzureChatOpenAI(
    api_key=AZURE_API_KEY,
    azure_endpoint=AZURE_ENDPOINT,
    azure_deployment=CHAT_DEPLOYMENT,
    openai_api_version=API_VERSION,
    temperature=0.3,
    max_tokens=800
)

async def main():
    session_id = str(uuid.uuid4())

    # MCP 서버 실행 파라미터
    server_params = StdioServerParameters(
        command="python",
        args=["mcp_server.py"], # MCP 를 서버 경로
    )

    async with stdio_client(server_params) as (read, write):
        async with ClientSession(read, write) as session:
            await session.initialize()
            tools = await load_mcp_tools(session)

            # React Agent 생성 (툴 목록 제공)
            agent = create_react_agent(model, tools)

            while True:
                user_input = input("\n질문을 입력해주세요: ").strip()
                if not user_input:
                    break

                # React Agent에게 메시지 전달
                messages = [
                    {"role": "user", "content": f"{user_input}\n(session_id: {session_id})"}
                ]
                result = await agent.ainvoke({"messages": messages})

                print("\n[에이전트 응답]")
                print(result)

                cont = input("\n계속하시겠습니까? (예/아니오): ").strip()
                if not cont.startswith("예"):
                    break

    if __name__ == "__main__":
        asyncio.run(main())
```

```
async def main():
    session_id = str(uuid.uuid4())

    # MCP 서버 실행 파라미터
    server_params = StdioServerParameters(
        command="python",
        args=["mcp_server.py"], # MCP 를 서버 경로
    )

    async with stdio_client(server_params) as (read, write):
        async with ClientSession(read, write) as session:
            await session.initialize()
            tools = await load_mcp_tools(session)

            # React Agent 생성 (툴 목록 제공)
            agent = create_react_agent(model, tools)

            while True:
                user_input = input("\n질문을 입력해주세요: ").strip()
                if not user_input:
                    break

                # React Agent에게 메시지 전달
                messages = [
                    {"role": "user", "content": f"{user_input}\n(session_id: {session_id})"}
                ]
                result = await agent.ainvoke({"messages": messages})

                print("\n[에이전트 응답]")
                print(result)

                cont = input("\n계속하시겠습니까? (예/아니오): ").strip()
                if not cont.startswith("예"):
                    break

    if __name__ == "__main__":
        asyncio.run(main())
```

## Section

### ReAct Agent 기반 실습

# 실행(1)

```
(py3_11_9) azureuser@b2b28-ML:~/work/jupyter/3_llm_service$ python mcp_react_client.py
[07/21/25 09:48:03] INFO    Processing request of type ListToolsRequest                                         server.py:625
질문을 입력해주세요 : ARS 시스템 점검은 언제 하나요 ?                                         [07/21/25 09:48:14] INFO    Processing request of type CallToolRequest                                         server.py:625
[07/21/25 09:48:15] INFO    HTTP Request: POST                                         _client.py:1025
[07/21/25 09:48:15] INFO    https://b2b28-md6zczaej-eastus2.cognitiveservices.azure.com/openai/deployments/text-embedding-3-small/embeddings?api-version=2024-12-01-preview "HTTP/1.1 200 OK"
[07/21/25 09:48:17] INFO    HTTP Request: POST                                         _client.py:1025
[07/21/25 09:48:17] INFO    https://b2b28-md6zczaej-eastus2.cognitiveservices.azure.com/openai/deployments/gpt-4.1-mini/chat/completions?api-version=2024-12-01-preview "HTTP/1.1 200 OK"

[에이전트 응답]
{'messages': [HumanMessage(content='ARS 시스템 점검은 언제 하나요 ?\n(session_id: 0ddf632d-ae60-4003-ac97-5cd6d3a7c348)', additional_kwargs={}, response_metadata={}, id='d3c3e4d1-f322-44ef-8d32-c16e2437b3b5'), AIMessage(content='', additional_kwargs={'tool_calls': [{'id': 'call_MlZpYEnFpAcSn21dPu8Gmljn', 'function': {'arguments': {'question': "ARS 시스템 점검 일정", "session_id": "0ddf632d-ae60-4003-ac97-5cd6d3a7c348"}, 'name': 'rag_tool', 'type': 'function'}}, 'refusal': None}, response_metadata={'token_usage': {'completion_tokens': 47, 'prompt_tokens': 108, 'total_tokens': 155, 'completion_tokens_details': {'accepted_prediction_tokens': 0, 'audio_tokens': 0, 'reasoning_tokens': 0, 'rejected_prediction_tokens': 0}, 'prompt_tokens_details': {'audio_tokens': 0, 'cached_tokens': 0}}, 'model_name': 'gpt-4.1-mini-2025-04-14', 'system_fingerprint': 'fp_178c8d546f', 'id': 'chatmpl-BvhPy4axunpjGNhaTPPFQ4qfYPfCJ', 'service_tier': None, 'prompt_filter_results': [{"prompt_index": 0, 'content_filter_results': {"hate": {"filtered": False, "severity": "safe"}, "self_harm": {"filtered": False, "severity": "safe"}, "sexual": {"filtered": False, "severity": "safe"}, "violence": {"filtered": False, "severity": "safe"}}}], 'finish_reason': 'tool_calls', 'logprobs': None, 'content_filter_results': {}}, 'id='run--cc3c2295-c405-42e4-94a1-d4c8d467fd55-0', tool_calls=[{'name': 'rag_tool', 'args': {'question': 'ARS 시스템 점검 일정', 'session_id': '0ddf632d-ae60-4003-ac97-5cd6d3a7c348'}, 'id': 'call_MlZpYEnFpAcSn21dPu8Gmljn', 'type': 'tool_call'}], usage_metadata={'input_tokens': 108, 'output_tokens': 47, 'total_tokens': 155, 'input_token_details': {'audio': 0, 'cache_read': 0, 'output_token_details': {'audio': 0, 'reasoning': 0}}}, ToolMessage(content='\n "answer": "ARS 시스템은 매월 점검합니다.",\n "references": [\n   "ARS 시스템은 매월 점검하며 장애 발생 시 즉시 전산팀에 보고합니다.",\n   "기지국 점검은 월 1회 정기적으로 수행되며 장애 이력은 3년간 보관합니다.",\n   "고객 접점에서 확인된 오류는 VOC 시스템을 통해 수집 및 분류됩니다."], name='rag_tool', id='2c8cc545-9f1d-4aec-860c-ced7e228a6a1', tool_call_id='call_MlZpYEnFpAcSn21dPu8Gmljn'), AIMessage(content='ARS 시스템 점검은 매월 정기적으로 진행됩니다. 점검 중 장애가 발생하면 즉시 전산팀에 보고하게 되어 있습니다. 추가로 궁금한 점 있으시면 말씀해 주세요.', additional_kwargs={'refusal': None}, response_metadata={'token_usage': {'completion_tokens': 45, 'prompt_tokens': 259, 'total_tokens': 304, 'completion_tokens_details': {'accepted_prediction_tokens': 0, 'audio_tokens': 0, 'reasoning_tokens': 0, 'rejected_prediction_tokens': 0}, 'prompt_tokens_details': {'audio_tokens': 0, 'cached_tokens': 0}}, 'model_name': 'gpt-4.1-mini-2025-04-14', 'system_fingerprint': 'fp_178c8d546f', 'id': 'chatmpl-BvhQ1FMPkfqWh8EwzNqp0PsZIIL6p', 'service_tier': None, 'prompt_filter_results': [{"prompt_index": 0, 'content_filter_results': {"hate": {"filtered": False, "severity": "safe"}, "self_harm": {"filtered": False, "severity": "safe"}, "sexual": {"filtered": False, "severity": "safe"}, "violence": {"filtered": False, "severity": "safe"}}}], 'finish_reason': 'stop', 'logprobs': None, 'content_filter_results': {"hate": {"filtered": False, "severity": "safe"}, "self_harm": {"filtered": False, "severity": "safe"}, "sexual": {"filtered": False, "severity": "safe"}, "violence": {"filtered": False, "severity": "safe"}}, 'id='run--4d859004-5c44-4446-9b3c-18db1001bf3d-0', usage_metadata={'input_tokens': 259, 'output_tokens': 45, 'total_tokens': 304, 'input_token_details': {'audio': 0, 'cache_read': 0, 'output_token_details': {'audio': 0, 'reasoning': 0}}})]
```

## Section

### ReAct Agent 기반 실습

# 실행(2)

```
계속 하시겠습니까? (예/아니오) 예
질문을 입력해주세요: 이 문장을 정중하게 바꿔줘. "이메일 빨리 보내세요."
[07/21/25 09:48:34] INFO Processing request of type CallToolRequest server.py:625
[07/21/25 09:48:36] INFO HTTP Request: POST _client.py:1025
[https://b2b28-md6zcjaj-eastus2.cognitiveservices.azure.com/openai/deployments/gpt-4.1-mini/chat/completions?api-version=2024-12-01-preview] "HTTP/1.1 200 OK"

[에이전트 응답]
{'messages': [HumanMessage(content='이 문장을 정중하게 바꿔줘. "이메일 빨리 보내세요."\n(session_id: 0ddf632d-ae60-4003-ac97-5cd6d3a7c348)', additional_kwargs={}, response_metadata={}, id='40d05aaaf-dac3-4cac-9540-a42f541aaf74'), AIMessage(content='', additional_kwargs={'tool_calls': [{'id': 'call_LJqwnc2BiW034tS4JVseiks7', 'function': {'arguments': {'text': '이메일 빨리 보내세요.'}, 'name': 'rephrase_tool', 'type': 'function'}], 'refusal': None}, response_metadata={'token_usage': {'completion_tokens': 22, 'prompt_tokens': 119, 'total_tokens': 141, 'completion_tokens_details': {'accepted_prediction_tokens': 0, 'audio_tokens': 0, 'reasoning_tokens': 0, 'rejected_prediction_tokens': 0}, 'prompt_tokens_details': {'audio_tokens': 0, 'cached_tokens': 0}}, 'model_name': 'gpt-4.1-mini-2025-04-14', 'system_fingerprint': 'fp_178c8d546f', 'id': 'chatcmpl-BvhQI2MBF9rHsIdgWCpG7i1ks4SDv', 'service_tier': None, 'prompt_filter_results': [{'prompt_index': 0, 'content_filter_results': {'hate': {'filtered': False, 'severity': 'safe'}, 'self_harm': {'filtered': False, 'severity': 'safe'}, 'sexual': {'filtered': False, 'severity': 'safe'}, 'violence': {'filtered': False, 'severity': 'safe'}}}], 'finish_reason': 'tool_calls', 'logprobs': None, 'content_filter_results': {}}, id='run--8f5adbea-69ad-40a1-8de6-3844e4f9e52c-0', tool_calls=[{'name': 'rephrase_tool', 'args': {'text': '이메일 빨리 보내세요.'}, 'id': 'call_LJqwnc2BiW034tS4JVseiks7', 'type': 'tool_call'}], usage_metadata={'input_tokens': 119, 'output_tokens': 22, 'total_tokens': 141, 'input_token_details': {'audio': 0, 'cache_read': 0}, 'output_token_details': {'audio': 0, 'reasoning': 0}}}, ToolMessage(content='이메일을 빠른 시일 내에 보내주시면 감사하겠습니다.', name='rephrase_tool', id='f4077f24-b32c-4d97-aa52-085f16a9ed10', tool_call_id='call_LJqwnc2BiW034tS4JVseiks7'), AIMessage(content='이메일을 빠른 시일 내에 보내주시면 감사하겠습니다.', additional_kwargs={'refusal': None}, response_metadata={'token_usage': {'completion_tokens': 18, 'prompt_tokens': 165, 'total_tokens': 183, 'completion_tokens_details': {'accepted_prediction_tokens': 0, 'audio_tokens': 0, 'reasoning_tokens': 0, 'rejected_prediction_tokens': 0}, 'prompt_tokens_details': {'audio_tokens': 0, 'cached_tokens': 0}}, 'model_name': 'gpt-4.1-mini-2025-04-14', 'system_fingerprint': 'fp_178c8d546f', 'id': 'chatcmpl-BvhQKN6vdw2xDto3aIRvJ1MeGMxTU', 'service_tier': None, 'prompt_filter_results': [{'prompt_index': 0, 'content_filter_results': {'hate': {'filtered': False, 'severity': 'safe'}, 'self_harm': {'filtered': False, 'severity': 'safe'}, 'sexual': {'filtered': False, 'severity': 'safe'}, 'violence': {'filtered': False, 'severity': 'safe'}}}], 'finish_reason': 'stop', 'logprobs': None, 'content_filter_results': {'hate': {'filtered': False, 'severity': 'safe'}, 'self_harm': {'filtered': False, 'severity': 'safe'}, 'sexual': {'filtered': False, 'severity': 'safe'}, 'violence': {'filtered': False, 'severity': 'safe'}}}, id='run--f9f39b55-0b47-4516-a1ad-14c199196065-0', usage_metadata={'input_tokens': 165, 'output_tokens': 18, 'input_token_details': {'audio': 0, 'cache_read': 0}, 'output_token_details': {'audio': 0, 'reasoning': 0}}])]

계속 하시겠습니까? (예/아니오): 아니오
(py3_11_9) azureuser@b2b28-ML:~/work/jupyter/3_llm_service$
```

# GEN AI 인텐시브 과정

Section 1. 실무 적용

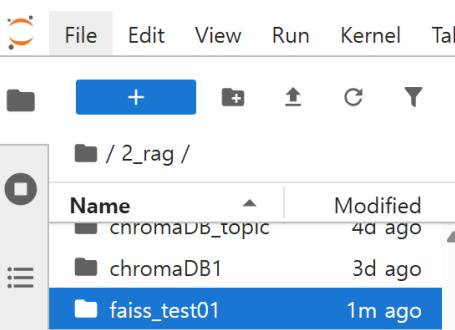
---

Section 1-2. 벡터 DB

## Section

### FAISS기초

# FAISS 기초



```
[2]: from langchain_huggingface import HuggingFaceEmbeddings
      from langchain_community.vectorstores import FAISS

[3]: # 임베딩 모델 설정
      embedding_model = HuggingFaceEmbeddings(model_name="sentence-transformers/paraphrase-multilingual-MiniLM-L12-v2")

[4]: # 저장할 텍스트 데이터
      texts = ["사과는 빨갛다", "바다는 파랗다", "개나리는 노랗다"]

[5]: # FAISS에 저장
      faiss_store = FAISS.from_texts(texts, embedding=embedding_model)

[6]: # 검색
      query = "딸기는 빨갛다"
      result = faiss_store.similarity_search(query, k=2)

[7]: # 검색 결과
      for doc in result:
          print(doc.page_content)

사과는 빨갛다
개나리는 노랗다

[8]: # 로컬 저장 & 불러오기
      faiss_store.save_local("./faiss_test01")

[10]: loaded_store = FAISS.load_local(
        folder_path = "./faiss_test01",
        embeddings = embedding_model,
        allow_dangerous_deserialization=True
    )
```

## Section

### Qdrant 설치

# Qdrant 버전 확인

<https://github.com/qdrant/qdrant/releases>

Apr 1  
general  
v1.13.6  
4db98ec  
Compare

## v1.13.6

### Change log

#### Improvements

- #6279 - In query API, read vectors/payloads once at shard level instead of in every segment, greatly improve search performance when there's lots of segments
- #6276 - In query API, don't send huge vectors/payloads over internal network, defer reads to greatly improve search performance
- #6260 - Improve performance of resharding transfers, make them faster on slow disks or with high memory pressure

#### Bug fixes

- #6259 - Fix point estimation in resharding transfers, showing a more reliable ETA
- #6233 - Fix order\_by not always including all values for a point if there are multiple

#### Assets

File	Size	Last Updated
qdrant-aarch64-apple-darwin.tar.gz	23.3 MB	Apr 1
qdrant-aarch64-unknown-linux-musl.tar.gz	24.6 MB	Apr 1
qdrant-x86_64-apple-darwin.tar.gz	25.1 MB	Apr 1

## Section

### Qdrant 설치

# Qdrant 다운로드 및 압축 해제

```
azureuser@b2b28-ML:~$ ls
work
azureuser@b2b28-ML:~$ cd work/
azureuser@b2b28-ML:~/work$ wget https://github.com/qdrant/qdrant/releases/download/v1.13.6/qdrant-x86_64-unknown-linux-gnu.tar.gz
Length: 27753850 (26M) [application/octet-stream]
Saving to: 'qdrant-x86_64-unknown-linux-gnu.tar.gz'

qdrant-x86_64-unknown-linux-gnu.tar.gz    100%[=====] 26.47M --.-KB/s   in 0.1s

2025-07-11 05:02:55 (250 MB/s) - 'qdrant-x86_64-unknown-linux-gnu.tar.gz' saved [27753850/27753850]

azureuser@b2b28-ML:~/work$ ls
data jupyter qdrant-x86_64-unknown-linux-gnu.tar.gz work
azureuser@b2b28-ML:~/work$ mkdir app
azureuser@b2b28-ML:~/work$ ls
app data jupyter qdrant-x86_64-unknown-linux-gnu.tar.gz work
azureuser@b2b28-ML:~/work$ mv qdrant-x86_64-unknown-linux-gnu.tar.gz app/
azureuser@b2b28-ML:~/work$ ls
app data jupyter work
azureuser@b2b28-ML:~/work$ cd app/
azureuser@b2b28-ML:~/work/app$ ls
qdrant-x86_64-unknown-linux-gnu.tar.gz
azureuser@b2b28-ML:~/work/app$ tar -xvzf qdrant-x86_64-unknown-linux-gnu.tar.gz
qdrant
azureuser@b2b28-ML:~/work/app$ ls
qdrant qdrant-x86_64-unknown-linux-gnu.tar.gz
```

## Section

### Qdrant 설치

# Qdrant 실행

```
azureuser@b2b28-ML:~/work/app$ ls  
qdrant  qdrant-x86_64-unknown-linux-gnu.tar.gz  
azureuser@b2b28-ML:~/work/app$ ./qdrant
```



```
Version: 1.13.6, build: 4db98ecd  
Access web UI at http://localhost:6333/dashboard
```

```
2025-07-11T05:09:38.008664Z  WARN qdrant::settings: Config file not found: config/config  
2025-07-11T05:09:38.008692Z  WARN qdrant::settings: Config file not found: config/development  
2025-07-11T05:09:38.008807Z  INFO storage::content_manager::consensus::persistent: Initializing new raft state at ./storage/raft_state.json  
2025-07-11T05:09:38.029684Z  INFO qdrant: Distributed mode disabled  
2025-07-11T05:09:38.029740Z  INFO qdrant: Telemetry reporting enabled, id: a51b10f5-6267-4018-84af-9588e891fec1  
2025-07-11T05:09:38.029781Z  INFO qdrant: Inference service is not configured.  
2025-07-11T05:09:38.031173Z  WARN qdrant::actix::web_ui: Static content folder for Web UI './static' does not exist  
2025-07-11T05:09:38.031417Z  INFO qdrant::actix: TLS disabled for REST API  
2025-07-11T05:09:38.031502Z  INFO qdrant::actix: Qdrant HTTP listening on 6333  
2025-07-11T05:09:38.031527Z  INFO actix_server::builder: Starting 3 workers  
2025-07-11T05:09:38.031539Z  INFO actix_server::server: Actix runtime found; starting in Actix runtime  
2025-07-11T05:09:38.037644Z  INFO qdrant::tonic: Qdrant gRPC listening on 6334  
2025-07-11T05:09:38.037672Z  INFO qdrant::tonic: TLS disabled for gRPC API
```



## Section

### Qdrant 설치

# 포트 열기

The screenshot shows the Azure portal interface for creating a new inbound security rule. The top navigation bar includes a search bar, Copilot, and user information.

**Left Panel:** Shows the 'b2b28-ML' resource group and its resources. The 'b2b28-ML' VM is selected. A message indicates a new version of the experience is being used.

**Middle Panel:** The 'b2b28-ML | 네트워크' blade is open, showing the '네트워크 설정' section. Under '네트워크 설정', there is a '네트워크 관리자' link.

**Right Panel:** The '인바운드 보안 규칙 추가' (Add Inbound Security Rule) dialog is open. The configuration is as follows:

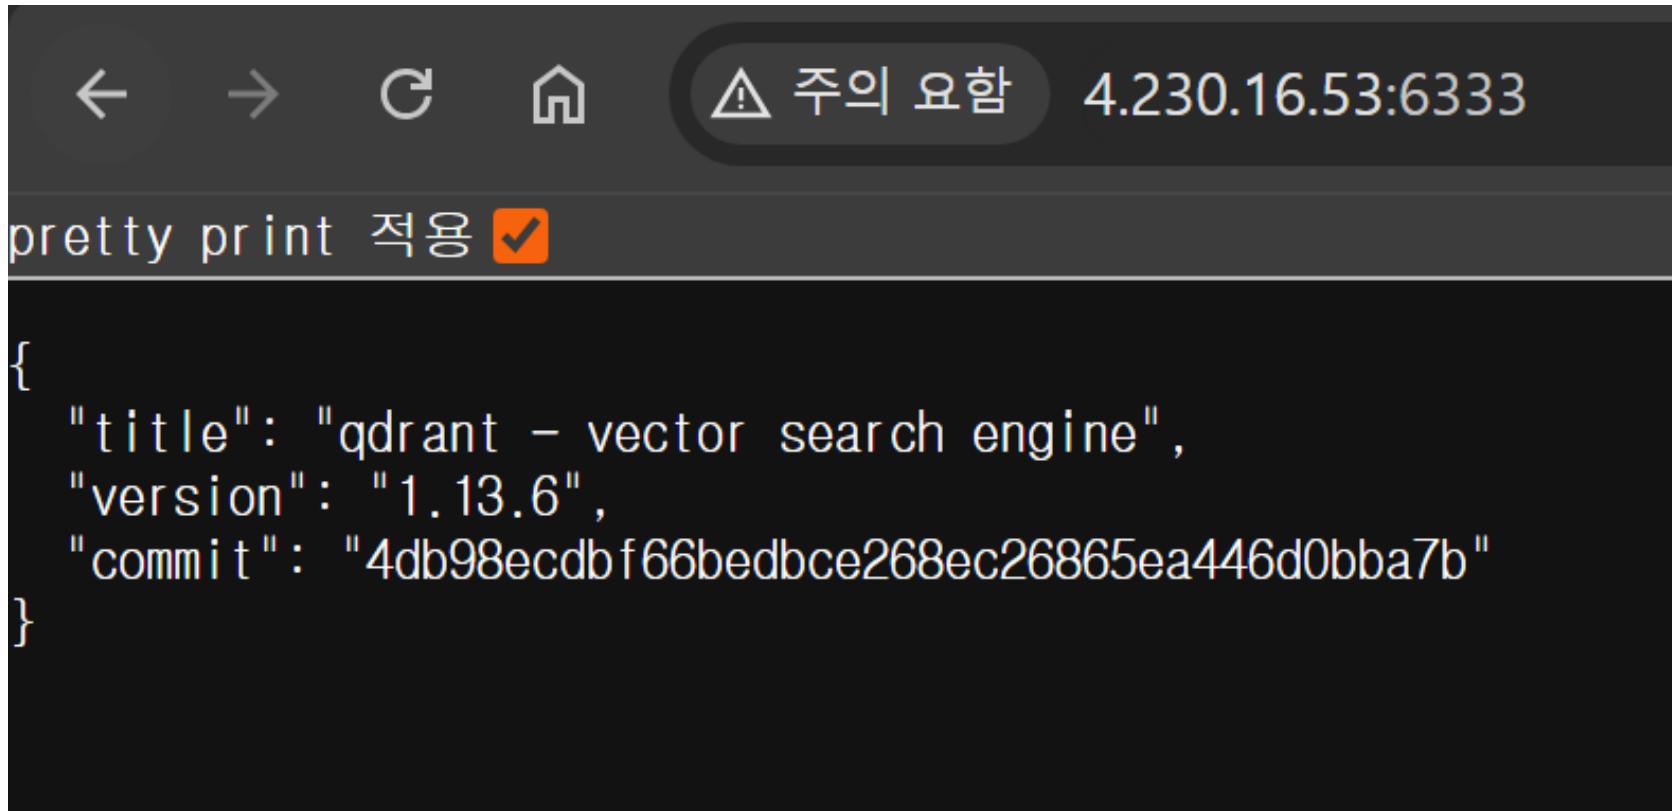
- 소스:** My IP address
- 원본 IP 주소/CIDR 범위:** 221.148.18.193
- 원본 포트 범위:** \*
- 대상 주소:** Any
- 서비스:** Custom
- 대상 포트 범위:** 6333
- 프로토콜:** Any (selected)
- 작업:** 허용 (Allow)

**Bottom Right:** Buttons for '추가' (Add) and '취소' (Cancel).

Section

Qdrant 설치

# 접속 확인



← → ⌛ ⌂ 주의 요함 4.230.16.53:6333

pretty print 적용 ✓

```
{  
  "title": "qdrant - vector search engine",  
  "version": "1.13.6",  
  "commit": "4db98ecdbf66bedbce268ec26865ea446d0bba7b"  
}
```

Section

Qdrant 설치

# Qdrant 클라이언트 설치

```
azureuser@b2b28-ML:~$ pyenv activate py3_11_9  
(py3_11_9) azureuser@b2b28-ML:~$ pip install qdrant-client  
(py3_11_9) azureuser@b2b28-ML:~$ pip install -U langchain-qdrant
```

## Section

### Qdrant 설치

# 라이브러리 불러오기 & 연결 & 임베딩 모델

```
from qdrant_client import QdrantClient
from qdrant_client.models import Distance, VectorParams
from langchain_qdrant import QdrantVectorStore
from langchain_huggingface import HuggingFaceEmbeddings
from langchain.schema import Document
```

# 1. Qdrant 클라이언트 연결

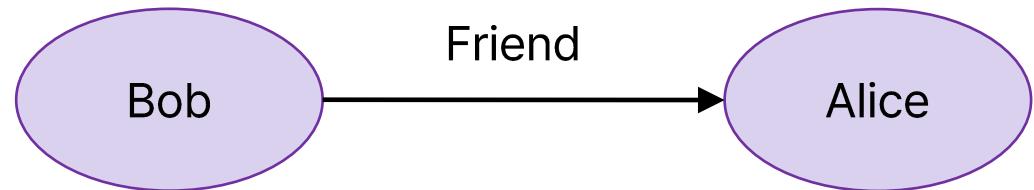
```
client = QdrantClient(host="localhost", port=6333)
```

# 2. 임베딩 모델

```
embedding_model = HuggingFaceEmbeddings(model_name="sentence-transformers/paraphrase-multilingual-MiniLM-L12-v2")
```

# GraphDB의 개념

- 데이터 간의 관계를 중점적으로 다루는 데이터베이스
  - 데이터를 노드(node)와 엣지(edge)라는 구조로 저장
    - 노드(node): 개체(ex: 사람, 도시, 제품 등)
    - 엣지(edge): 관계(ex: 친구, 가족, 위치, 연결 등)
    - 속성(property): 노드나 엣지의 추가정보



```
{  
  "node": {  
    "name": "Alice",  
    "type": "Person"  
  },  
  "edge": {  
    "type": "FRIEND",  
    "since": 2015  
  },  
  "target": {  
    "name": "Bob",  
    "type": "Person"  
  }  
}
```

# GraphDB의 필요성

- RDB(관계형 데이터베이스)의 한계
  - 테이블 조인(join)이 너무 많아져서 느림
  - 관계가 깊을수록 증가하는 복잡도
  - 유연한 스키마 설계 어려움
- GraphDB의 장점
  - 관계를 직접 표현 가능
  - 빠른 탐색 속도
  - 동적으로 확장 가능한 구조

## Section

### Neo4j 설치

# 기본 사용법

The screenshot shows the Neo4j Browser interface. At the top, there are navigation icons (back, forward, search, etc.) and a status bar indicating the URL is 4.230.16.53:7474/browser/ and a warning about a known issue. Below the status bar, a message box displays a Neo4j Cypher query:

```
neo4j$ CREATE (a:Person {name: "Alice"})-[:KNOWS]→(b:Person {name: "Bob"}) RETURN a, b;
```

Below the query, a blue banner provides information about product usage collection:

To help make Neo4j Browser better we collect information on product usage. Review your [settings](#) at any time.

In the bottom left corner, there is a terminal-like input field with the command:

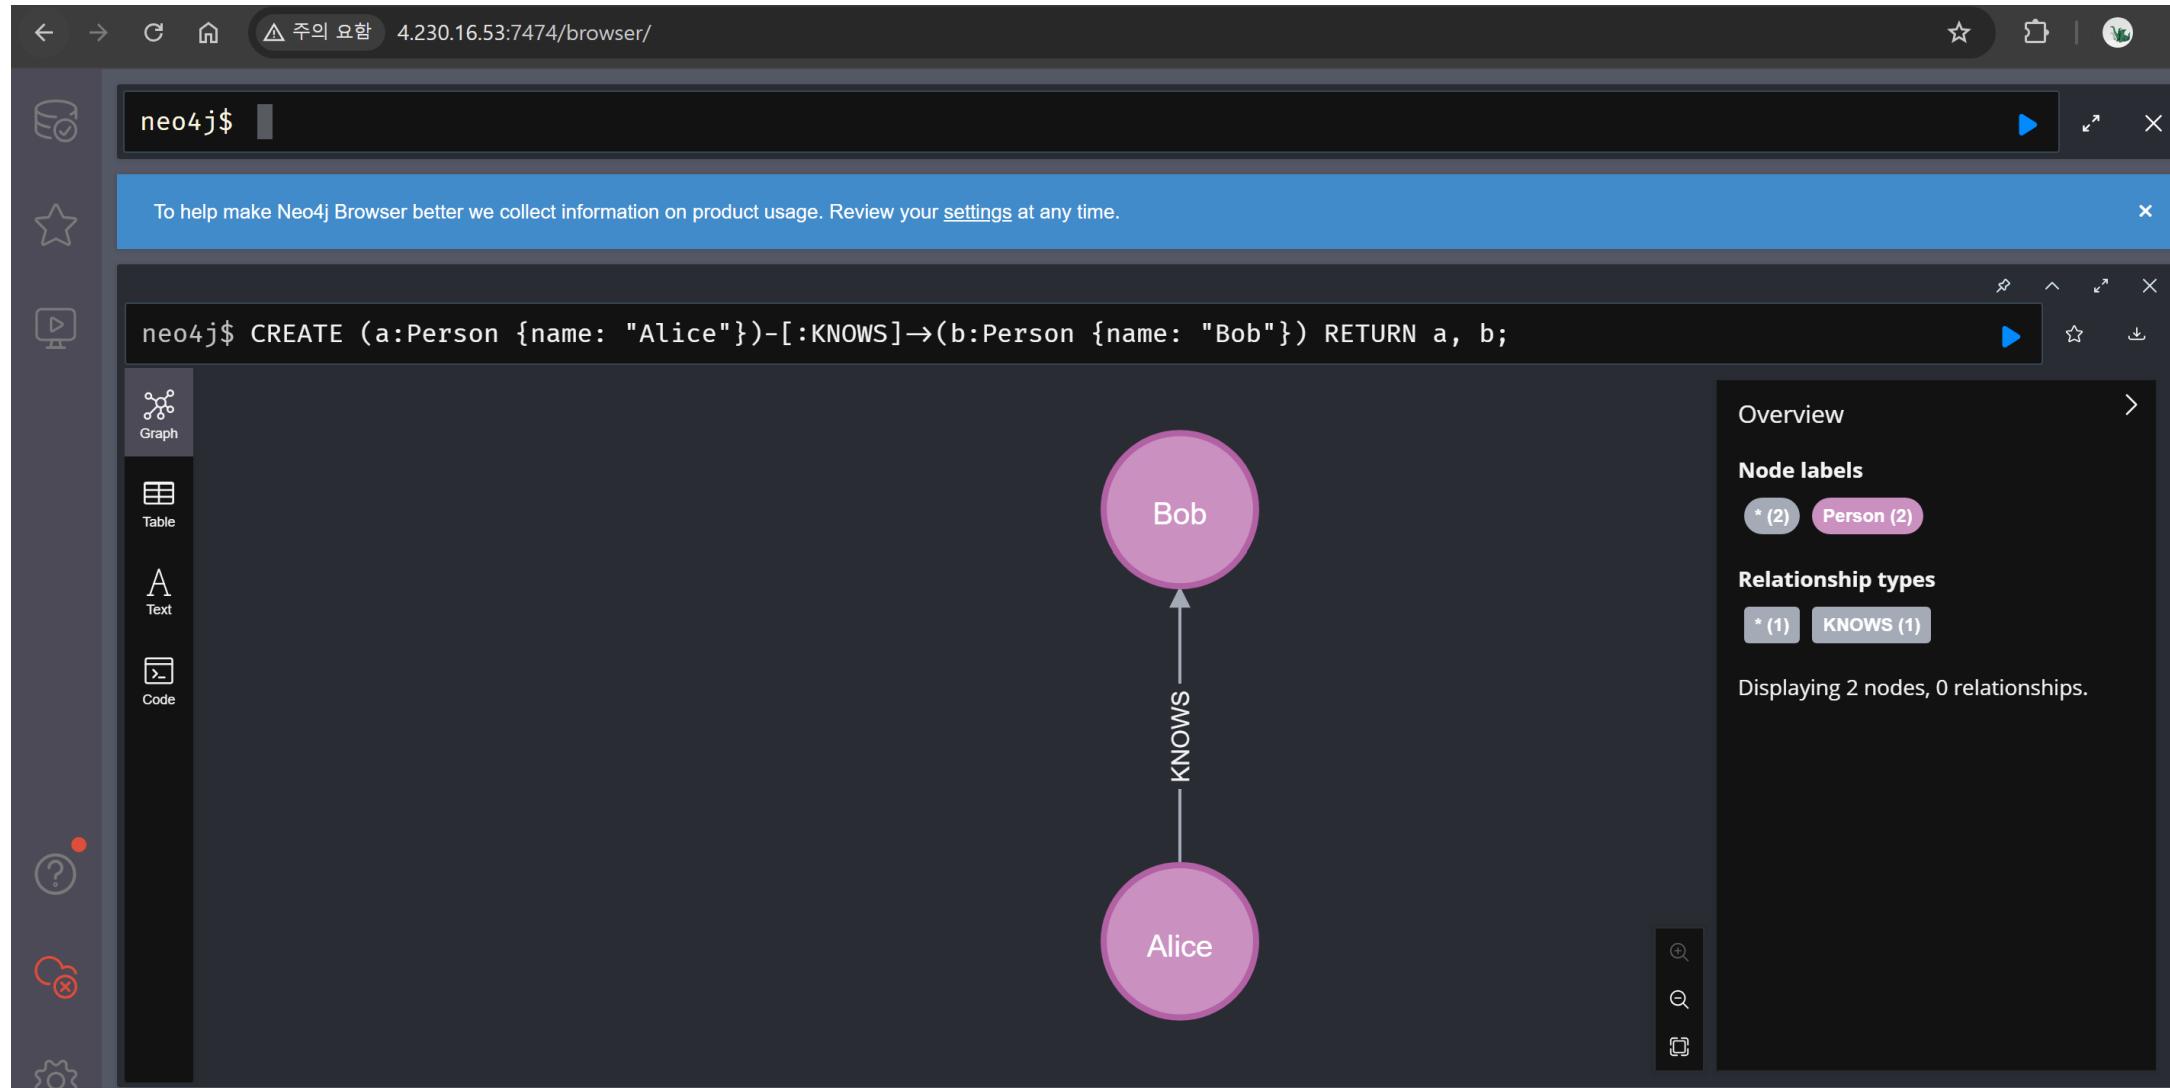
```
$ :play start
```

The Neo4j logo is visible at the bottom center of the browser window.

## Section

### Neo4j 설치

# 기본 사용법



## Section

### Neo4j 설치

# 기본 사용법

The screenshot shows the Neo4j Browser interface. The URL bar indicates the connection is to 4.230.16.53:7474/browser/. The main query entered is:

```
neo4j$ CREATE (a:Person {name: "Alice"})-[:KNOWS]→(b:Person {name: "Bob"}) RETURN a, b;
```

The results are displayed in two columns, labeled 'a' and 'b', showing the nodes created by the query. The 'Table' tab is selected in the sidebar.

a	b
<pre>{     "identity": 0,     "labels": [         "Person"     ],     "properties": {         "name": "Alice"     },     "elementId": "4:08bcc21f-d86b-410a-a3f9-1fe5ffe33b9a:0" }</pre>	<pre>{     "identity": 1,     "labels": [         "Person"     ],     "properties": {         "name": "Bob"     },     "elementId": "4:08bcc21f-d86b-410a-a3f9-1fe5ffe33b9a:1" }</pre>

At the bottom, a message states: "Added 2 labels, created 2 nodes, set 2 properties, created 1 relationship, started streaming 1 records after 15 ms and completed after 68 ms."

# GEN AI 인텐시브 과정

Section 1. 실무 적용

---

**Section 1-3. 실제 서비스**

## Section

### FastAPI 기초 사용법

# pydantic 라이브러리



클라이언트

json 형태로  
HTTP 전달

### HTTP 요청

POST /items HTTP/1.1  
Host: example.com  
Content-Type: application/json  
Content-Length: 61  
User-Agent: CustomClient/1.0

```
{  
    "name": "Book",  
    "price": 9.99,  
    "description": "novel"  
}
```

### 서버

HTTP 요청에서  
body에서  
JSON 추출

```
{  
    "name": "Book",  
    "price": 9.99,  
    "description": "novel"  
}
```

```
from fastapi import FastAPI  
from pydantic import BaseModel  
  
class Item(BaseModel):  
    name: str  
    price: float  
    description: str = None
```

추출한 JSON을  
Pydantic 객체로 변환

```
item = Item(**json_data)
```

## Section

### Nginx의 개념

# ASGI vs WSGI

ASGI는 비동기 지원



사용자

The Nginx logo in green text.

웹 서버



ASGI

The FastAPI logo with a lightning bolt icon.

웹 애플리케이션



사용자

The Nginx logo in green text.

웹 서버

A green unicorn icon representing WSGI.

WSGI



Flask

웹 애플리케이션



사용자

The Nginx logo in green text.

웹 서버

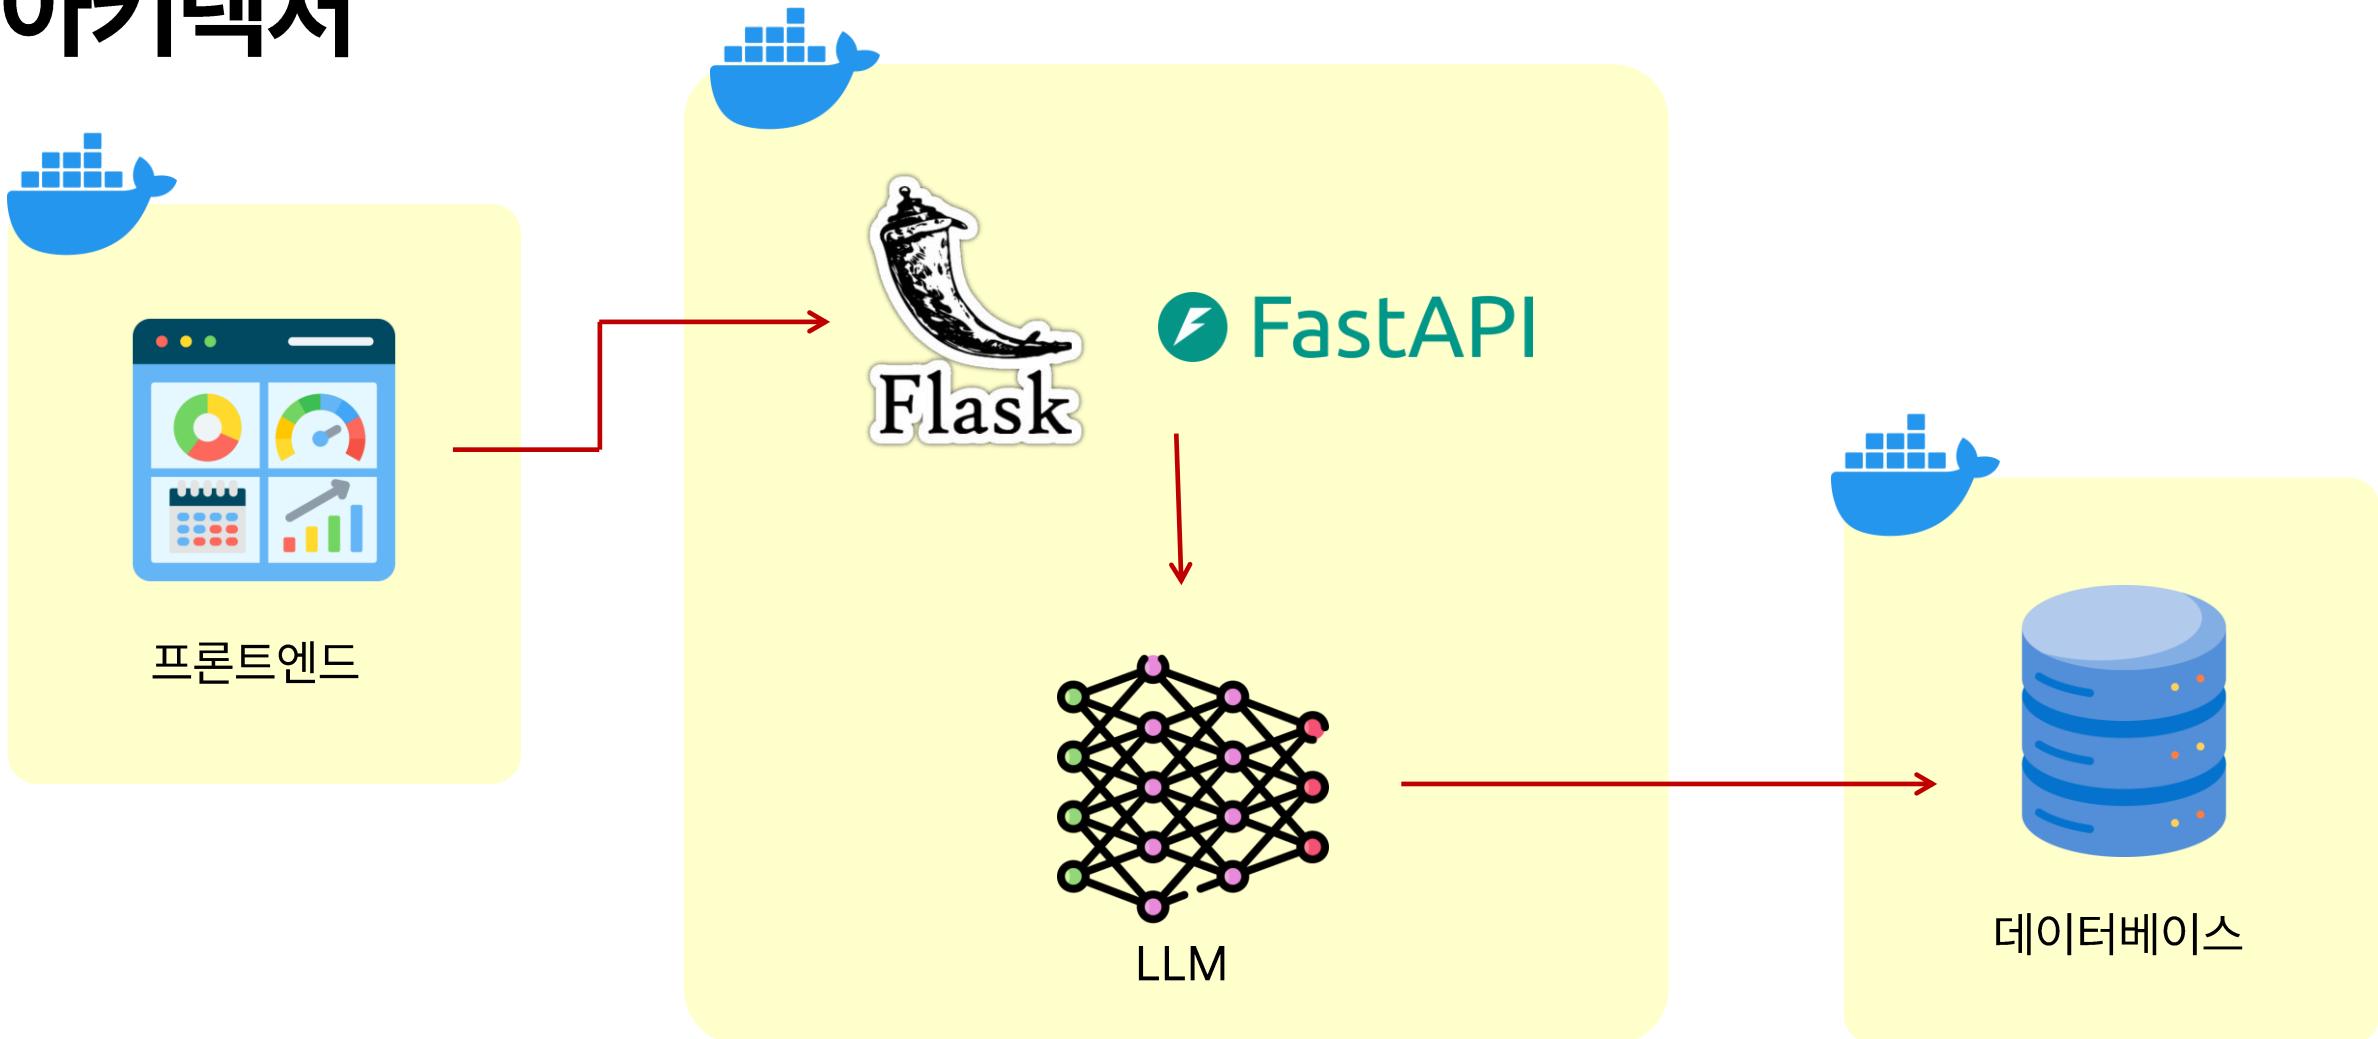
A green unicorn icon representing WSGI.

WSGI

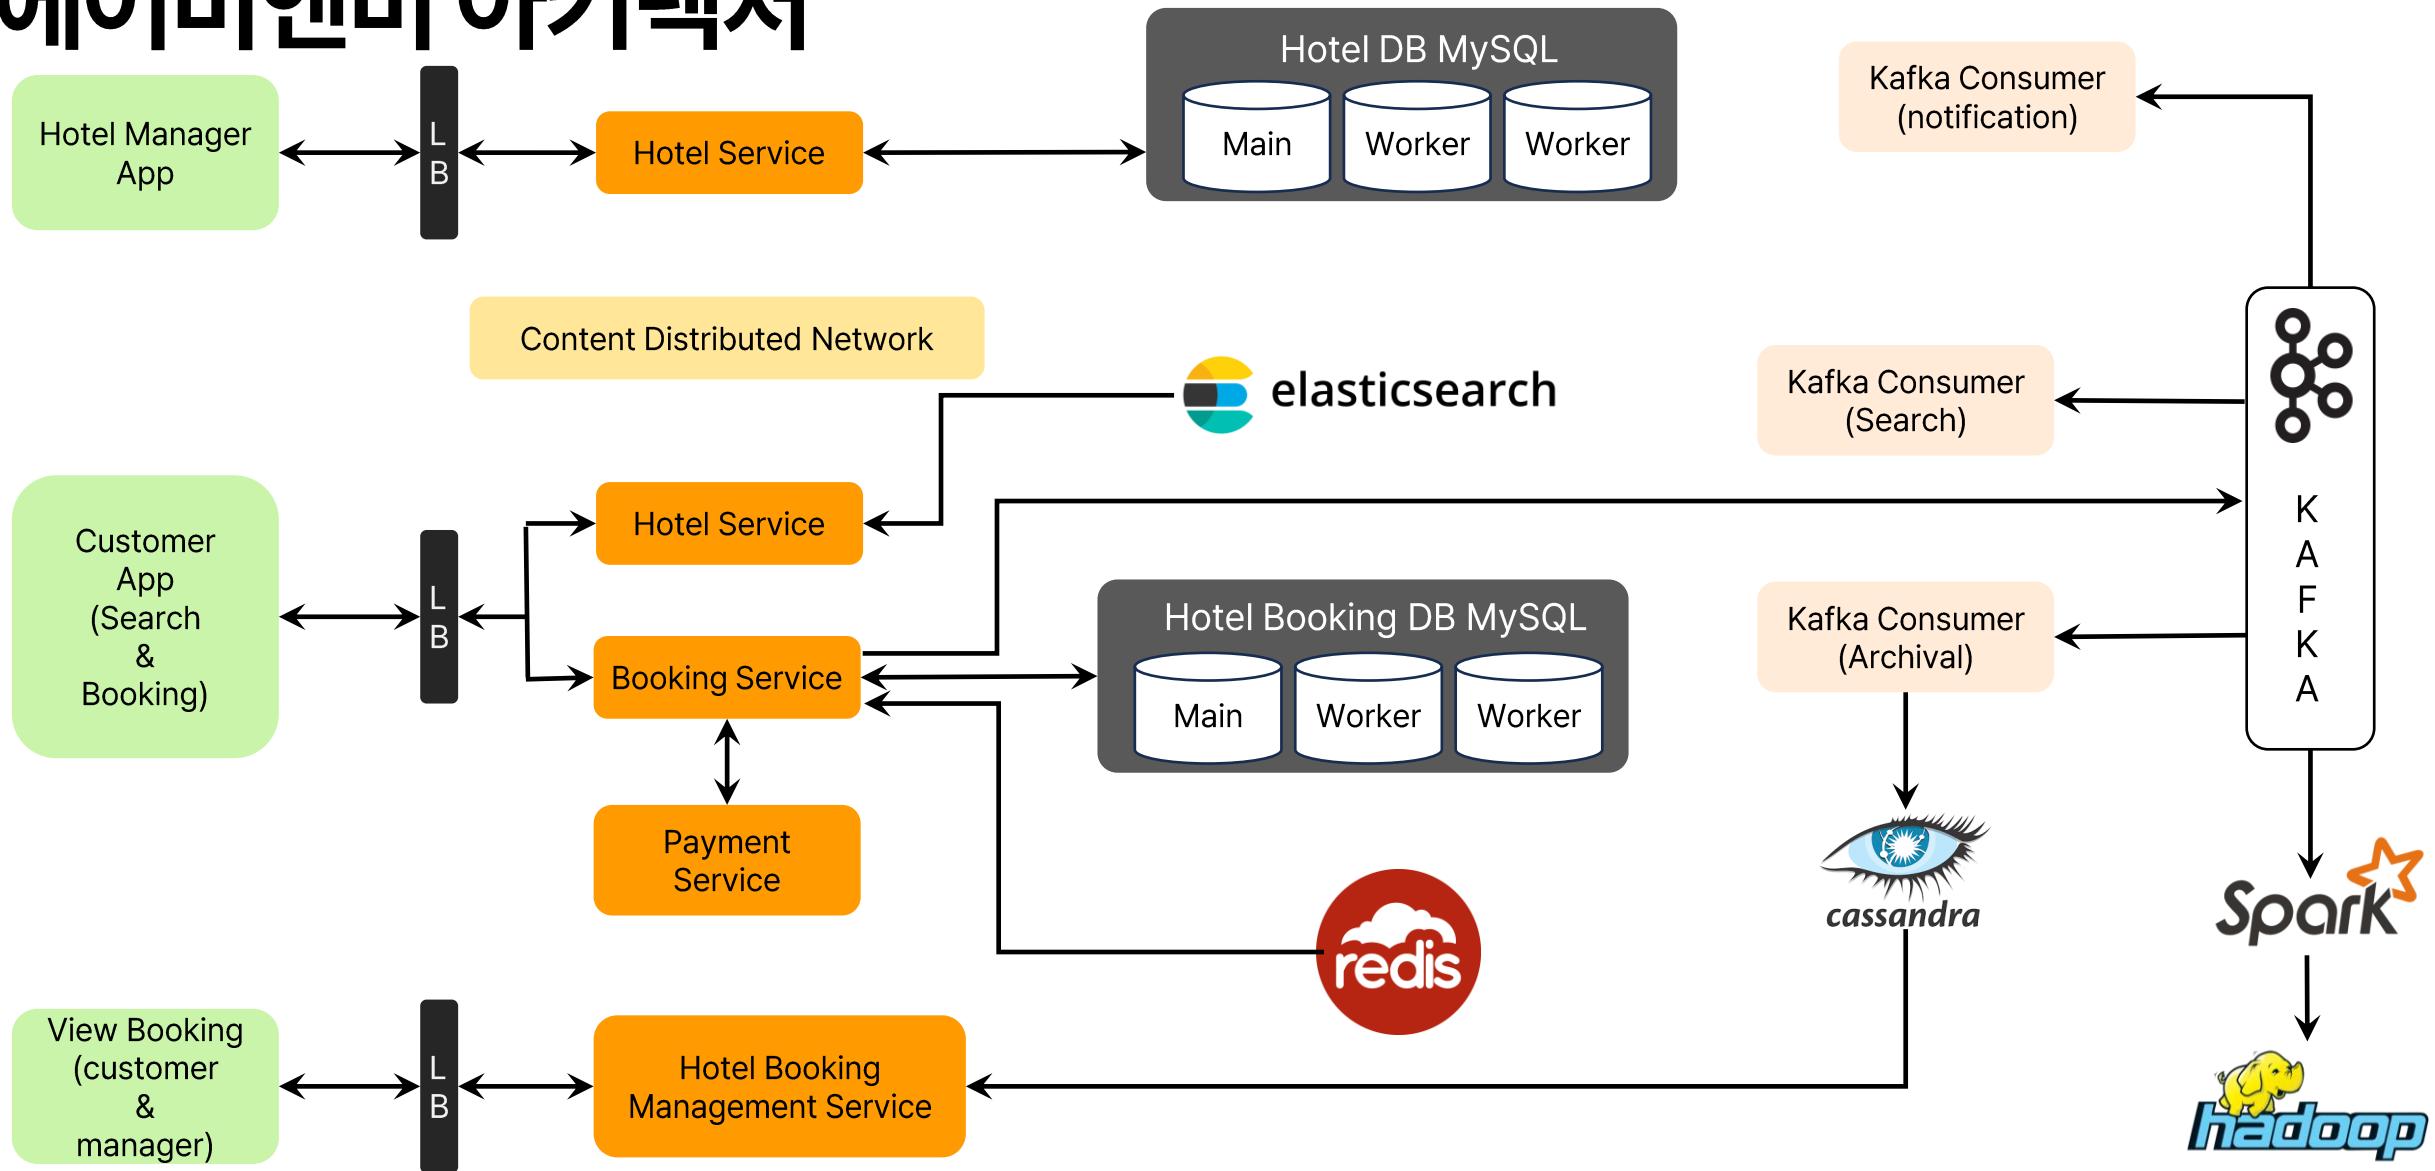
The Django logo in white text on a dark green background.

웹 애플리케이션

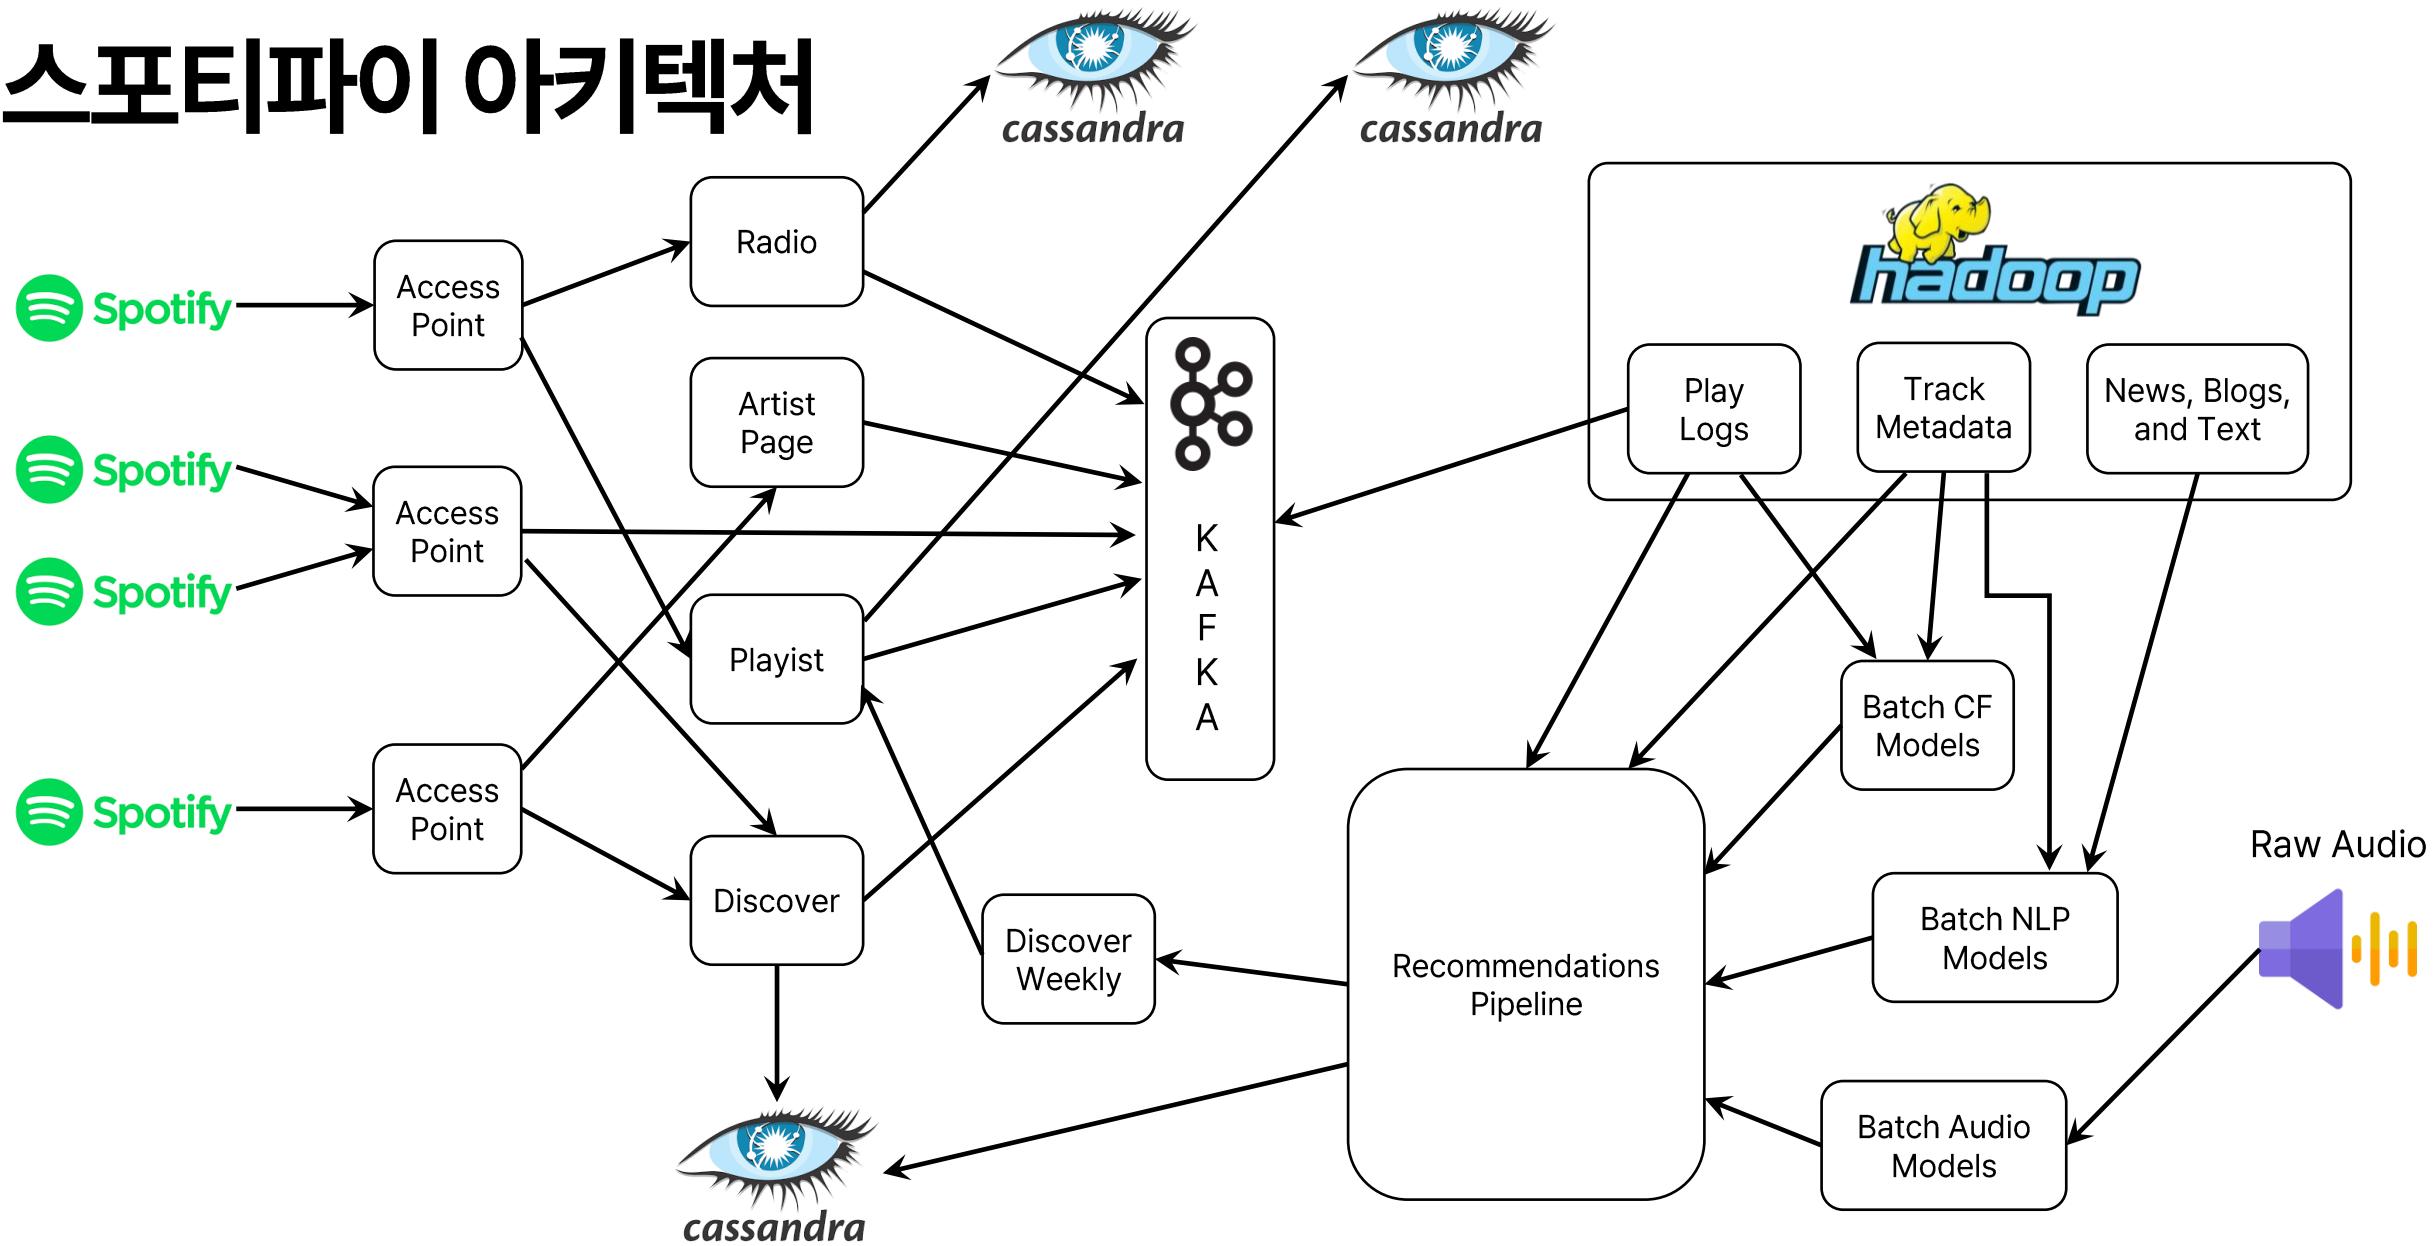
# 아키텍처



# 에어비앤비 아키텍처



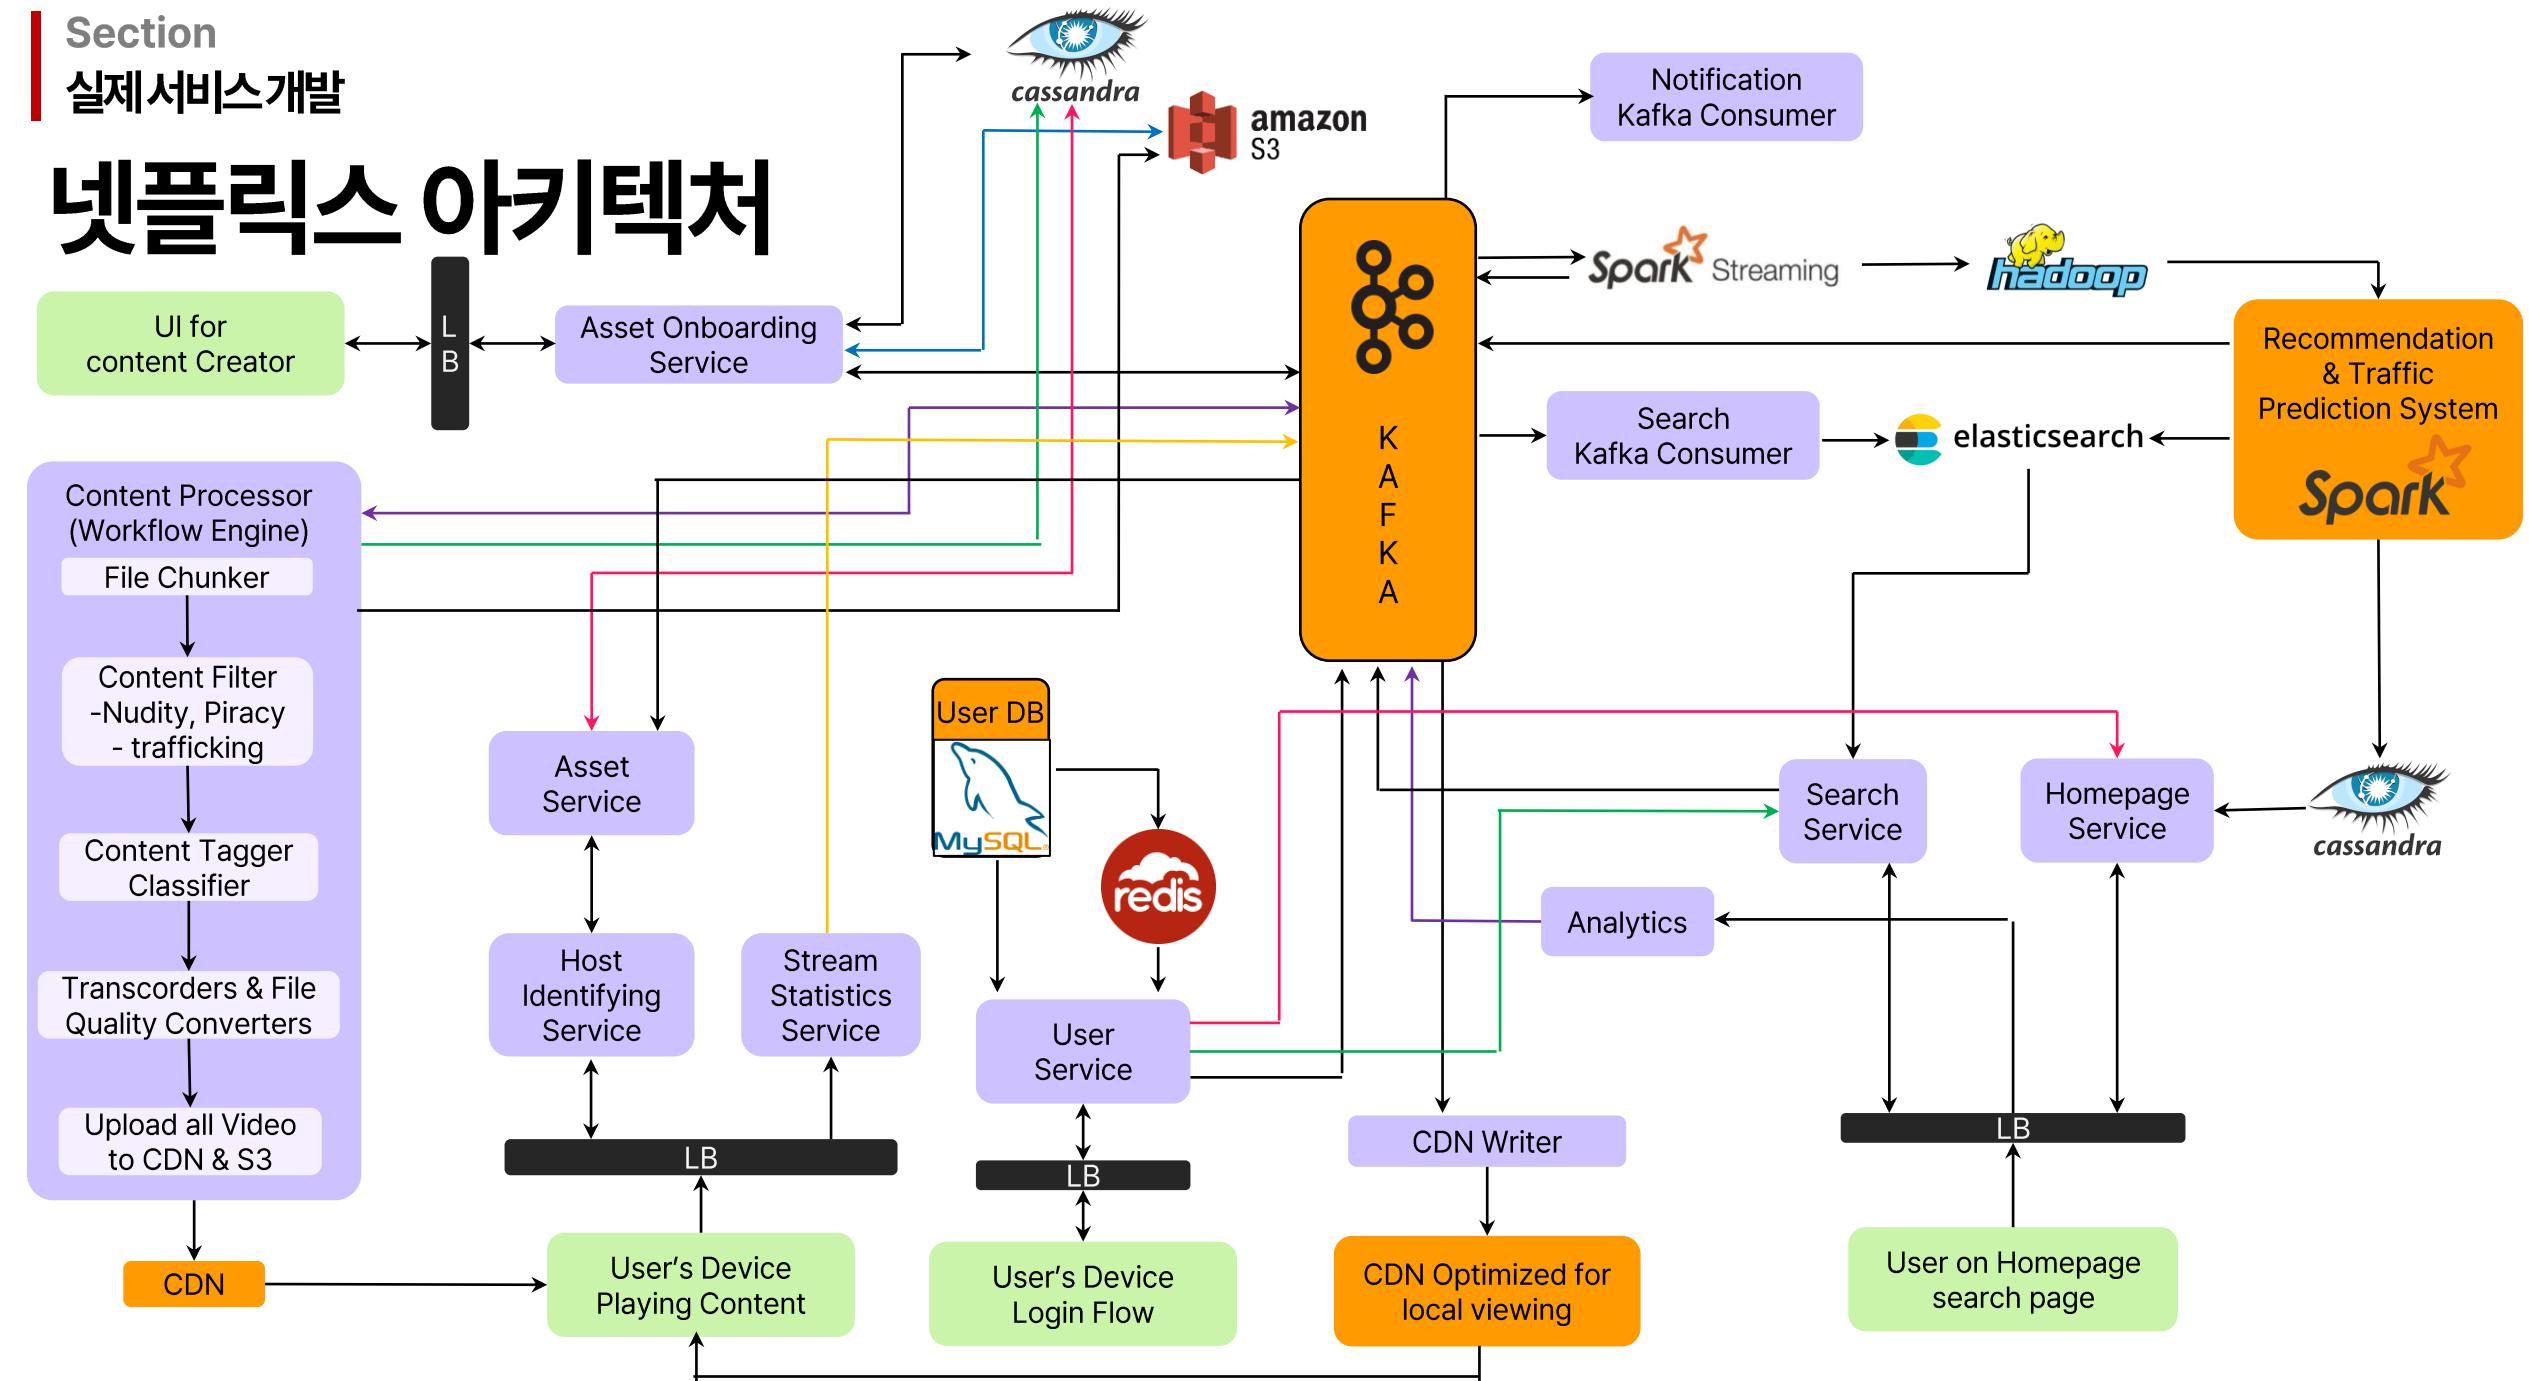
# 스포티파이 아키텍처



## Section

### 실제 서비스 개발

# 넷플릭스 아키텍처



감사합니다.

# Q & A